- Openshift
  Orchestrierungsservice zur  Bereitstellung, Verwaltung und Skalierung von Container-Anwendungen

- Deklaratives System
  Status wird in Resourcen (YAML/JSON) definiert und durch Controller hergestellt
  IaC – Infrastructure as  Code (https://blog.nelhage.com/post/declarative-configuration-management)

```
$ oc api-resources -o name --sort-by=name
alertmanagers.monitoring.coreos.com
apiservers.config.openshift.io
apiservices.apiregistration.k8s.io
appliedclusterresourcequotas.quota.openshift.io
authentications.config.openshift.io
authentications.operator.openshift.io
baremetalhosts.metal3.io
bindings
brokertemplateinstances.template.openshift.io
buildconfigs.build.openshift.io
builds.build.openshift.io
builds.config.openshift.io
catalogsources.operators.coreos.com
certificatesigningrequests.certificates.k8s.io
cloudcredentials.operator.openshift.io
clusterautoscalers.autoscaling.openshift.io
clusternetworks.network.openshift.io
clusteroperators.config.openshift.io
...
```

*Pod*
*Replicaset*
*Deployment*
*Service (svc)*
*Route*
*PersistentVolumeClaim*
*Secrets*
*Configmaps*

*Imagestream*
*BuildConfig*

*Node*
*PersistentVolume*

*Operator*
*CustomResourceDefinition*

openshift resources

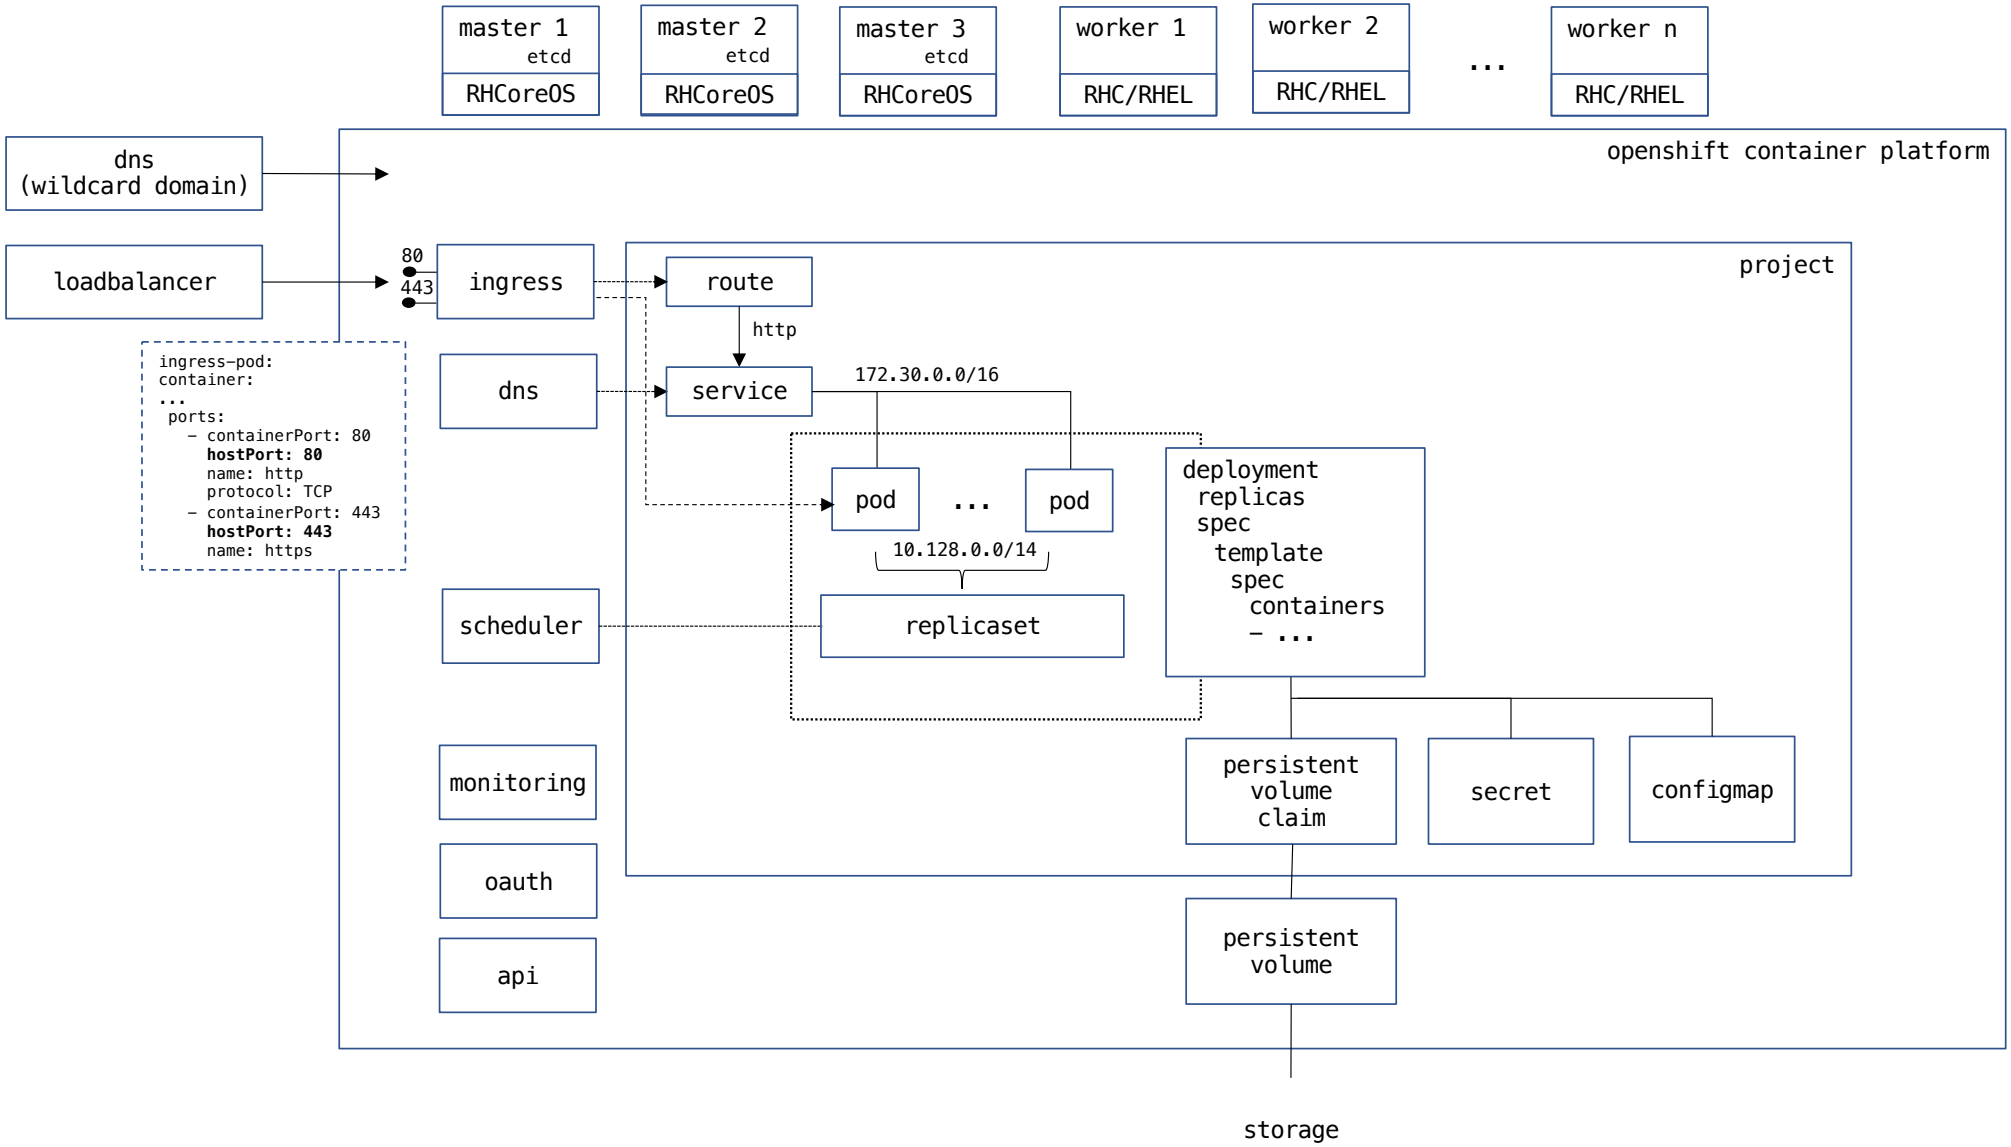
master 1
etcd
RHCoreOS
master 2
etcd
RHCoreOS
master 3
etcd
RHCoreOS
worker 1
RHC/RHEL
worker 2
RHC/RHEL
...
worker n
RHC/RHEL
dns
(wildcard domain)
loadbalancer
openshift container platform
80
443
ingress
route
http
project
ingress-pod:
container:
...
ports:
- containerPort: 80
hostPort: 80
name: http
protocol: TCP
- containerPort: 443
hostPort: 443
name: https
dns
service
172.30.0.0/16
pod
...
pod
10.128.0.0/14
deployment
replicas
spec
template
spec
containers
- ...
scheduler
replicaset
monitoring
oauth
api
persistent
volume
claim
secret
configmap
persistent
volume
storage

openshift

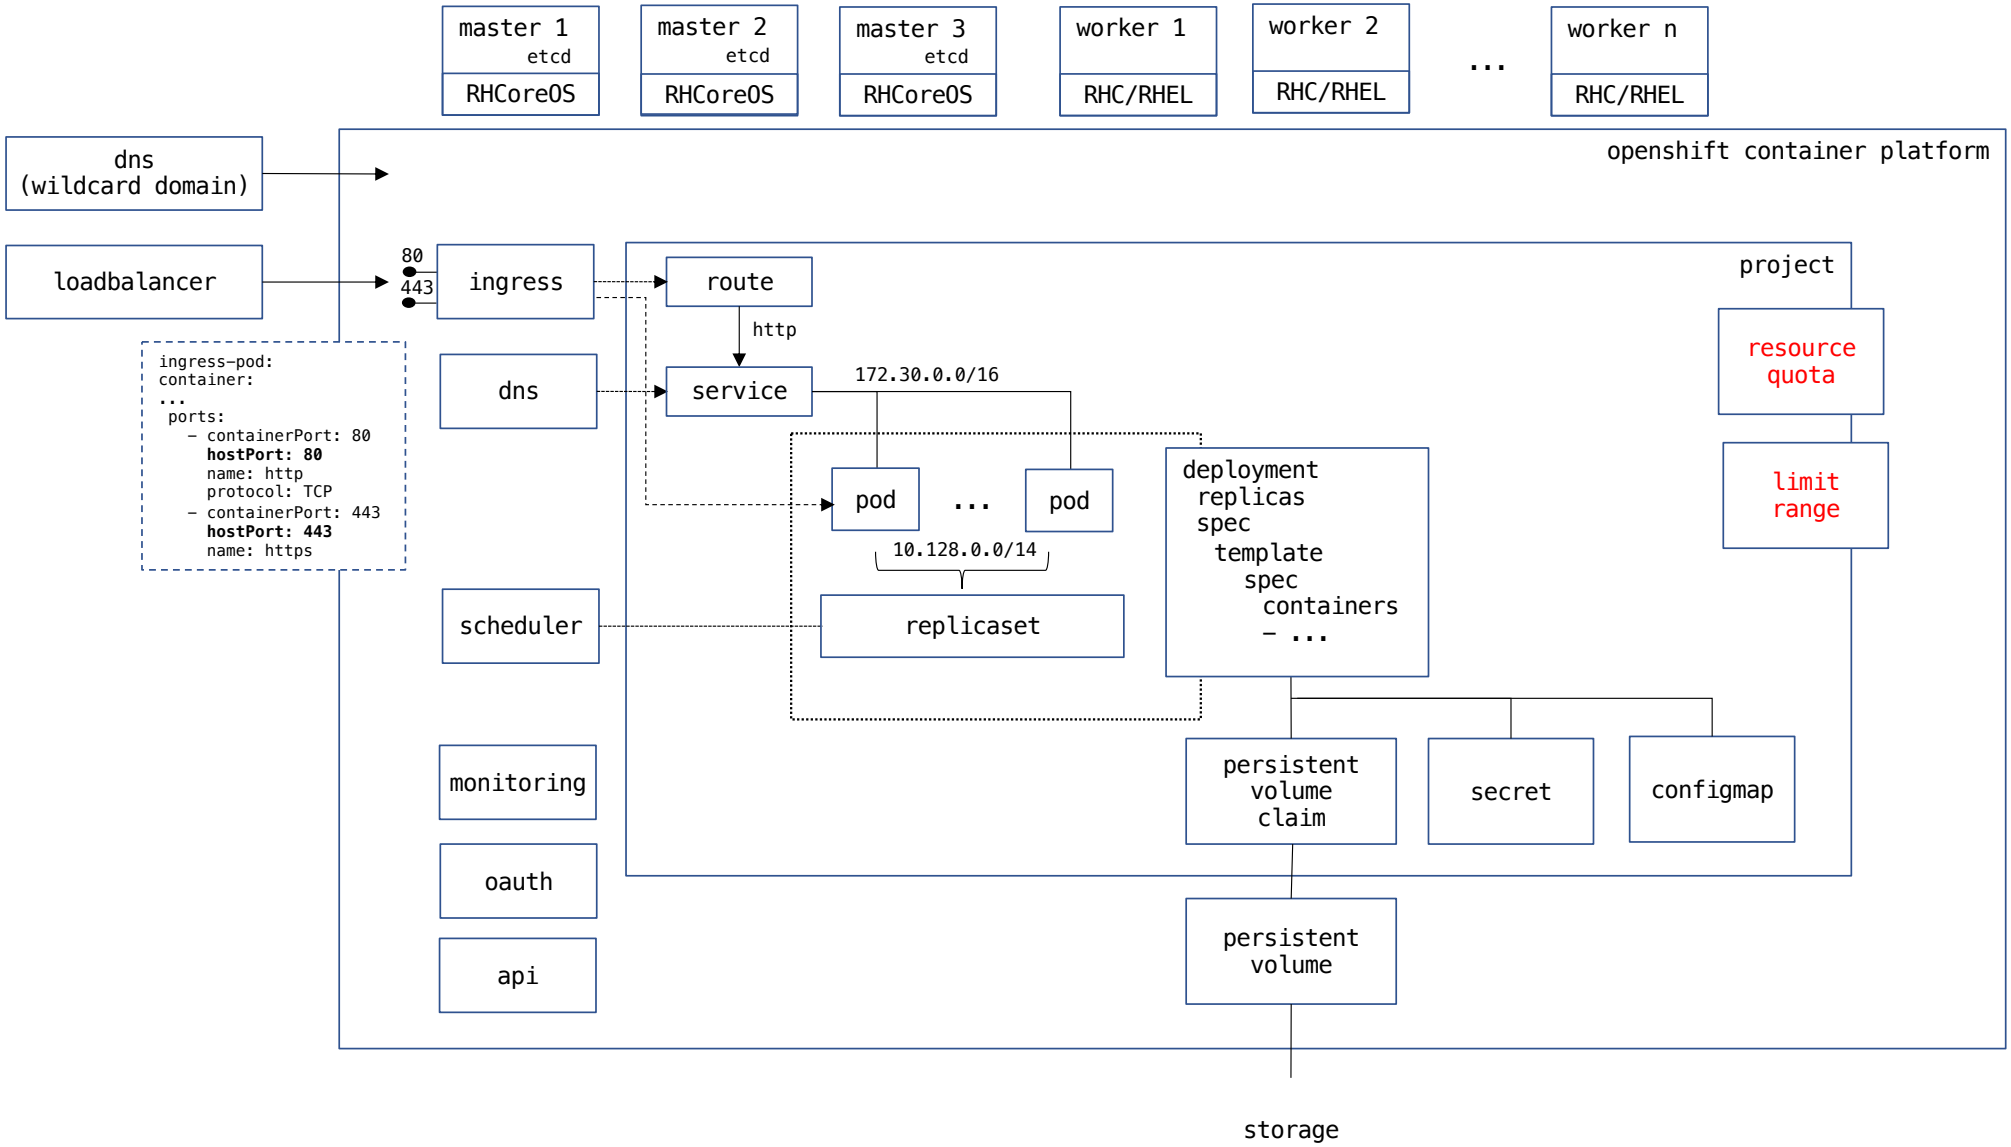
master 1
etcd
RHCoreOS
master 2
etcd
RHCoreOS
master 3
etcd
RHCoreOS
worker 1
RHC/RHEL
worker 2
RHC/RHEL
...
worker n
RHC/RHEL
dns
(wildcard domain)
loadbalancer
openshift container platform
80
443
ingress
route
http
project
resource
quota
ingress-pod:
container:
...
ports:
- containerPort: 80
hostPort: 80
name: http
protocol: TCP
- containerPort: 443
hostPort: 443
name: https
dns
service
172.30.0.0/16
pod
...
pod
10.128.0.0/14
deployment
replicas
spec
template
spec
containers
- ...
limit
range
scheduler
replicaset
monitoring
persistent
volume
claim
secret
configmap
oauth
persistent
volume
api
storage

openshift

# Openshift Resources (Manifest)

```
apiVersion: v1
kind: < Resource Type >
metadata:
  name: <name>
  namespace: <namespace>
  annotations:
    ...
  labels:
    app: <application-name>
    ...
spec:
  ...
  selector:
    <key>: <value>
  ...
status:
  ...
```

oc apply

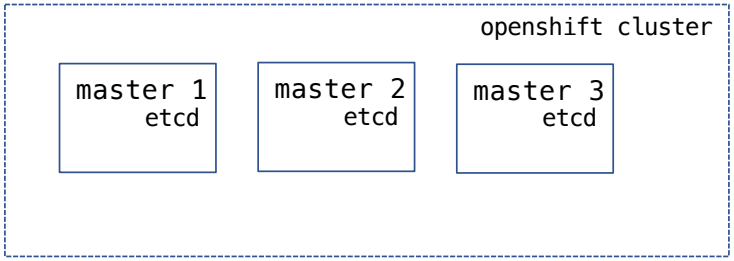


```
apiVersion: v1
kind: Pod
metadata:
  name: webserver
  namespace: do180
  labels:
    app: webserver
spec:
  containers:
  - image: quay.io/danielstraub/webserver:do180
    imagePullPolicy: Always
    ports:
    - containerPort: 8080
      protocol: TCP
  ...
```

```yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: wildfly
  labels:
    app: wildfly
spec:
  replicas: 2
  selector:
    matchLabels:
      app: wildfly
  template:
    metadata:
      labels:
        app: wildfly
    spec:
      containers:
      - name: wildfly
        image: quay.io/do288/wildfly:latest
        ports:
        - containerPort: 8080
          protocol: TCP
```

```yaml
apiVersion: v1
kind: Service
metadata:
  name: wildfly
  labels:
    app: wildfly
spec:
  type: ClusterIP
  selector:
    app: wildfly
  ports:
  - name: http
    protocol: TCP
    port: 8080
    targetPort: 8080
```

```yaml
apiVersion: route.openshift.io/v1
kind: Route
metadata:
  name: wildfly
  labels:
    app: wildfly
spec:
  host: sample.apps.eu46.prod.nextcle.com
  to:
    kind: Service
    name: wildfly
  tls:
    termination: edge
```

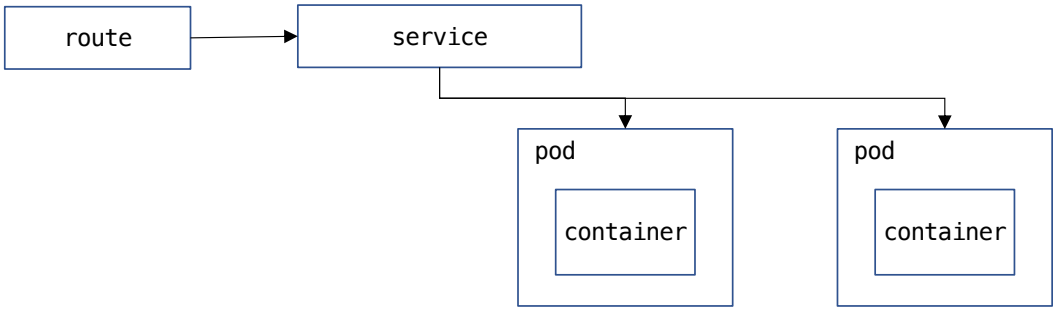


application

## Declarative :

```
$ ls
deployment.yml route.yml service.yml

$ oc apply -f .
deployment.apps/wildfly created
route.route.openshift.io/wildfly created
service/wildfly created
```

## Imperative :

```
$ oc new-app <container-image | git-repository>
--> Found container image 9a9e908 (9 days old) from quay.io for "quay.io/do288/wildfly"

    * An image stream tag will be created as "wildfly:latest" that will track this image

--> Creating resources ...
    imagestream.image.openshift.io "wildfly" created
    deployment.apps "wildfly" created
    service "wildfly" created
--> Success
```

- oc login –u <user> -p <password> <api-server-url>

- oc new-project <name>

- oc create –f <resource-yml>
  oc apply –f <resource-yml>

- oc status

- oc get <resource-type> [ <resource-name> ]

  - oc get pods
  - oc get deployment
  - oc get svc <service>
  - oc get events

- oc describe <resource-type> <resource-name>

- oc expose svc <service-name>

- oc logs <podname>

- oc exec –it <podname> -- <program>
- oc rsh <podname>
- oc cp <pod>:<locatio> <location>

- oc port-forward <podname> <local-port>:<remote-port>

- oc new-app <☺anything☺>

- oc delete <resource-type> <resource-name>

- oc rollout latest deployment <deployment-name>

https://docs.openshift.com/container-platform/4.15/cli_reference/openshift_cli/developer-cli-commands.html

```
$ oc new-app --help
Create a new application by specifying source code, templates, and/or images

...

Usage:
  oc new-app (IMAGE | CONTAINTERFILE | SOURCE | TEMPLATE | ...) [flags]


Beispiele:

$ oc new-app quay.io/do288/nginx --name ngnix
```

Container-Image

```
$ oc new-app php:7.3~https://github.com/.../php-hello
```

Builder-Image (s2i)

Git-Projekt (Source)

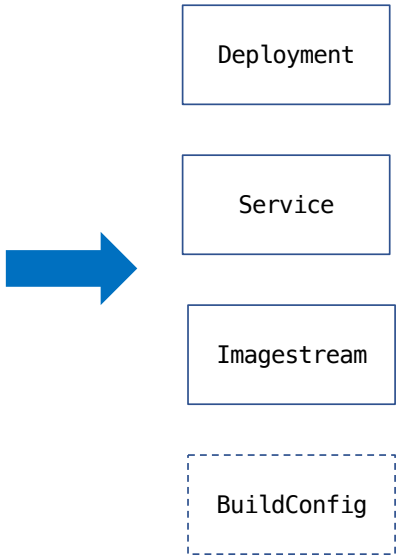


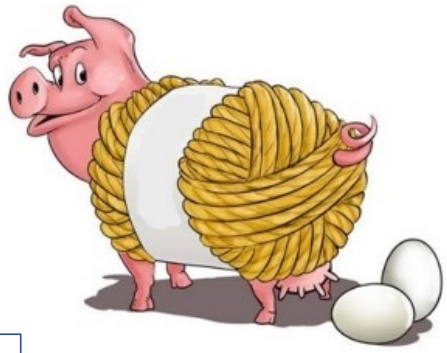

```
$ oc create deployment --image=quay.io/danielstraub/toolbox -o yaml toolbox -- bash -c 'sleep infitity'
apiVersion: apps/v1
kind: Deployment
metadata:
  name: toolbox
  labels:
    app: toolbox
spec:
  replicas: 1
  selector:
    matchLabels:
      app: toolbox
  template:
    metadata:
    labels:
        app: toolbox
    spec:
      containers:
      - command:
        - bash
        - -c
        - sleep infitity
        image: quay.io/danielstraub/toolbox
        name: toolbox
```

```
$ oc create service clusterip webserver --tcp=80:8080 -o yaml
apiVersion: v1
kind: Service
metadata:
  name: webserver
  labels:
    app: webserver
spec:
  ports:
  - name: 80-8080
    port: 80
    protocol: TCP
    targetPort: 8080
  selector:
    app: webserver
  type: ClusterIP

$ oc create route edge --hostname do180.<wildcard-doamin> --service webserver --insecure-policy=Redirect webserver -o yaml
apiVersion: route.openshift.io/v1
kind: Route
metadata:
  name: webserver
  labels:
    app: webserver
spec:
  host: do180.apps.eu410.prod.nextcle.com
  port:
    targetPort: http
  tls:
    insecureEdgeTerminationPolicy: Redirect
    termination: edge
  to:
    name: webserver
```

```
oc create route -help

Available Commands:
  edge        Create a route that uses edge TLS termination
  passthrough Create a route that uses passthrough TLS termination
  reencrypt   Create a route that uses reencrypt TLS termination
```

## odo innerLoop:

```
odo dev →  build run
odo dev –debug → build debug

verwenden eines 'nicht-Default' Commando:

odo dev --buildCommand=... --runCommand=
```

```
workstation:demo $ odo dev [ --no-watch ] [ --platform podman ]

 __
/  \__     Developing using the "java-springboot" Devfile
\__/  \    Namespace: architecture-setup
/  \__/    odo version: v3.15.0 (10b5e8a8f)
\__/

↪ Running on the cluster in Dev mode
•  Waiting for Kubernetes resources  ...
 ✓  Added storage m2 to component
===================
⚠  Pod is Pending
===================
 ✓  Pod is Running
 ✓  Syncing files into the container [801ms]
 ✓  Building your application in container (command: build) [1m]
 •  Executing the application (command: run)  ...

↪ Dev mode
Web console accessible at http://localhost:20000/

Keyboard Commands:
[Ctrl+c] - Exit and delete resources from the cluster
     [p] - Manually apply local changes to the application on the cluster
 ✓  Finished executing the application (command: run) [7s]
```

## odo outerLoop:

```
odo deploy →  deploy  (composite-Commando)
```

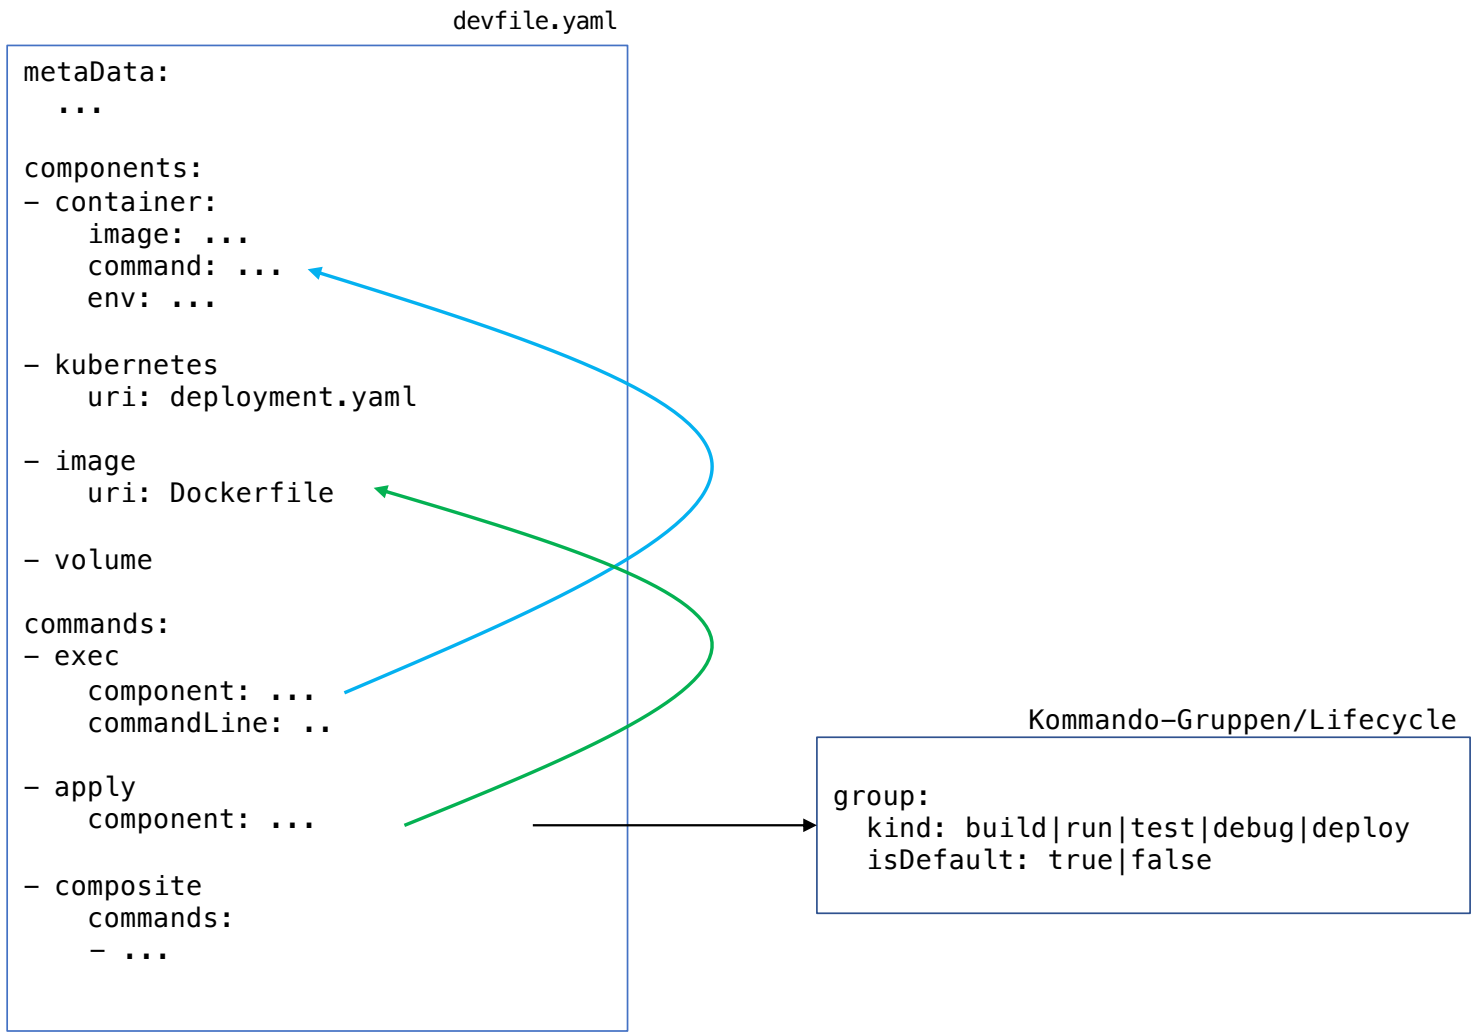
devfile.yaml
metaData:
...
components:
- container:
image: ...
command: ...
env: ...
- kubernetes
uri: deployment.yaml
- image
uri: Dockerfile
- volume
commands:
- exec
component: ...
commandLine: ..
- apply
component: ...
- composite
commands:
- ...
Kommando-Gruppen/Lifecycle
group:
kind: build|run|test|debug|deploy
isDefault: true|false

```
/
├── bin
├── etc
├── lib
├── opt
├── proc
├── usr
└── var
    └── lib
        └── container
            └── storage

/
systemd─┬─agetty
        ├─auditd───{auditd}
        ├─chronyd
        ├─crond
        ├─dbus-daemon
        ├─irqbalance
        ├─ksmtuned───sleep
        ├─polkitd───6*[{polkitd}]
        ├─rpc.idmapd
        ├─rpc.mountd
        ├─rpc.statd
        ├─rpcbind
        ├─rsyslogd───2*[{rsyslogd}]
        ├─sshd─── sshd───bash
        ├─systemd-journal
        ├─systemd-logind
        ├─systemd-machine
        ├─systemd-udevd
        └─ < app >
```

```
Container: namespace mnt, pid, net, user, uts, ipc

  /
  ├── bin
  ├── etc
  ├── lib
  ├── opt
  ├── proc
  ├── usr
  └── var

  /
  └─ < app >
```

ip: 10.88.0.x

cgroups : mem, cpu

selinux : context security

secomp  : syscalls

eth0

cni0: 10.88.0.1

KERNEL

rootfull container

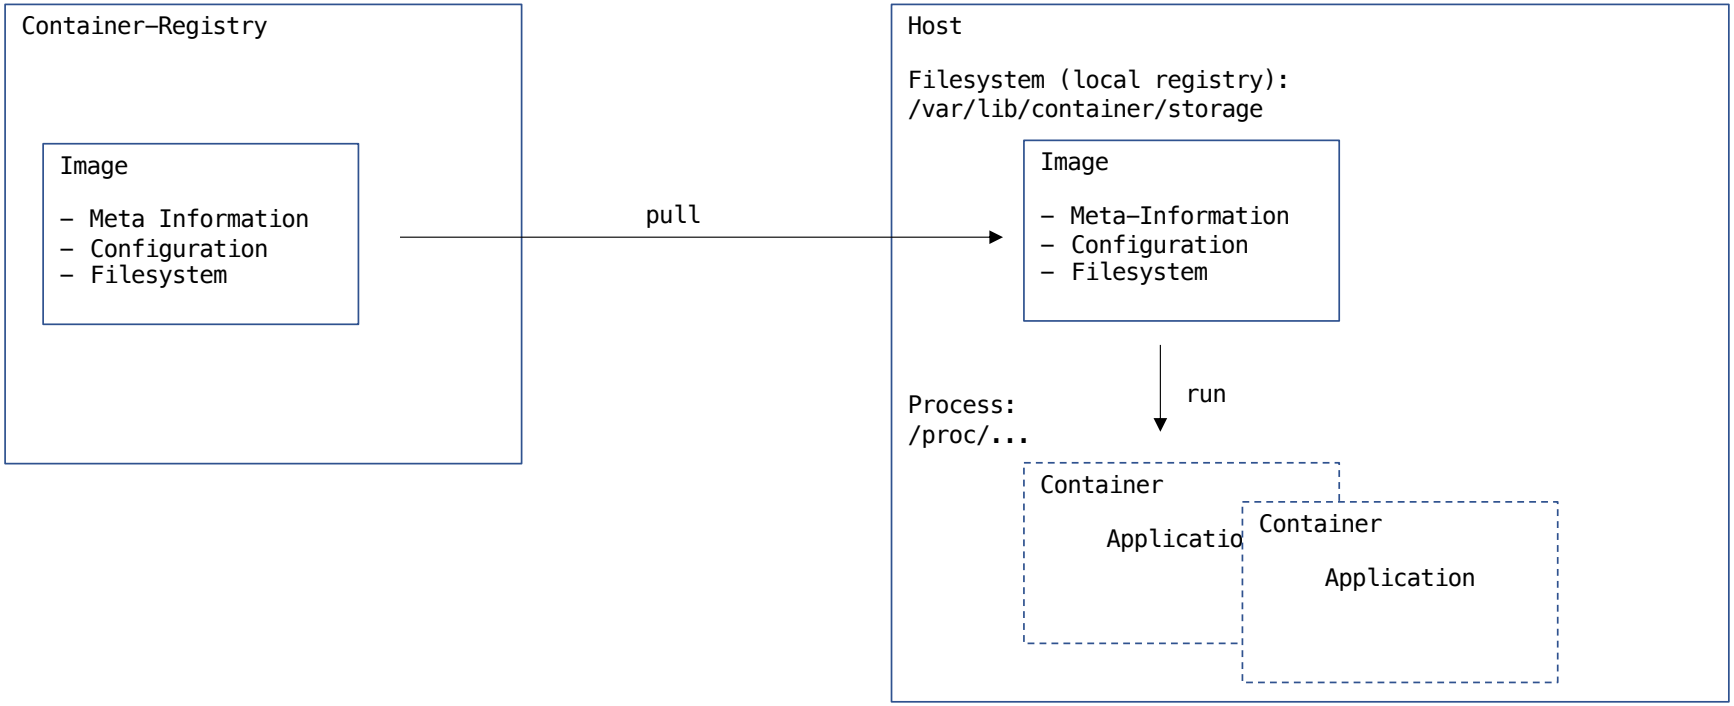


image and container

podman build - Containerfile

```
FROM ubi9/ubi
LABEL version=.... maintainer=
MAINTAINER daniel
ENV key=value
ARG version

ADD  http://repos/app-$version.tar /opt/app/
COPY webapp.war /opt/tomcat/webapps
RUN yum install -y tomcat  && \
         useradd tomcat && \
         chown -R 1000:1000 /opt/tomcat && \
         yum cleanup && rm /tmp/*


USER 1000
WORKDIR /opt/tomcat
VOLUMES /opt/tomcat/logs
EXPOSE [ 8080, 8001 ]
ENTRYPOINT [ "bin/sh -c" ]
CMD [ "bin/catalina.sh", "start" ]
```

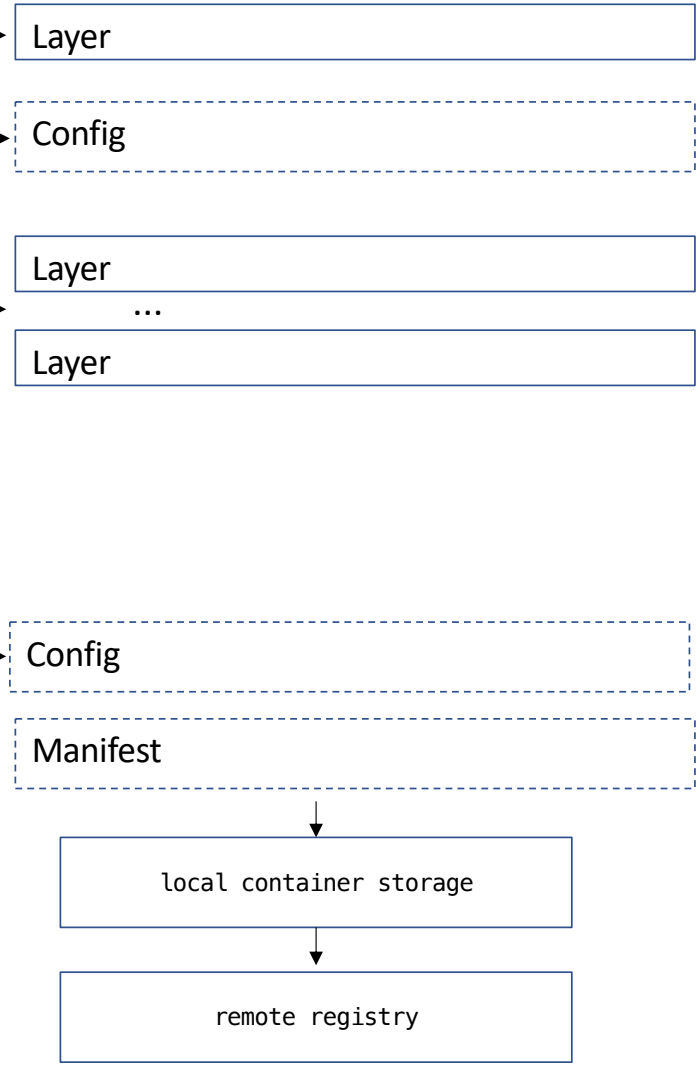


```
podman build -t my-tomcat .  <- Build-Dir


podman push <registry>/<name>:<tag>
```

build image

Beispiel: Image – Vererbung

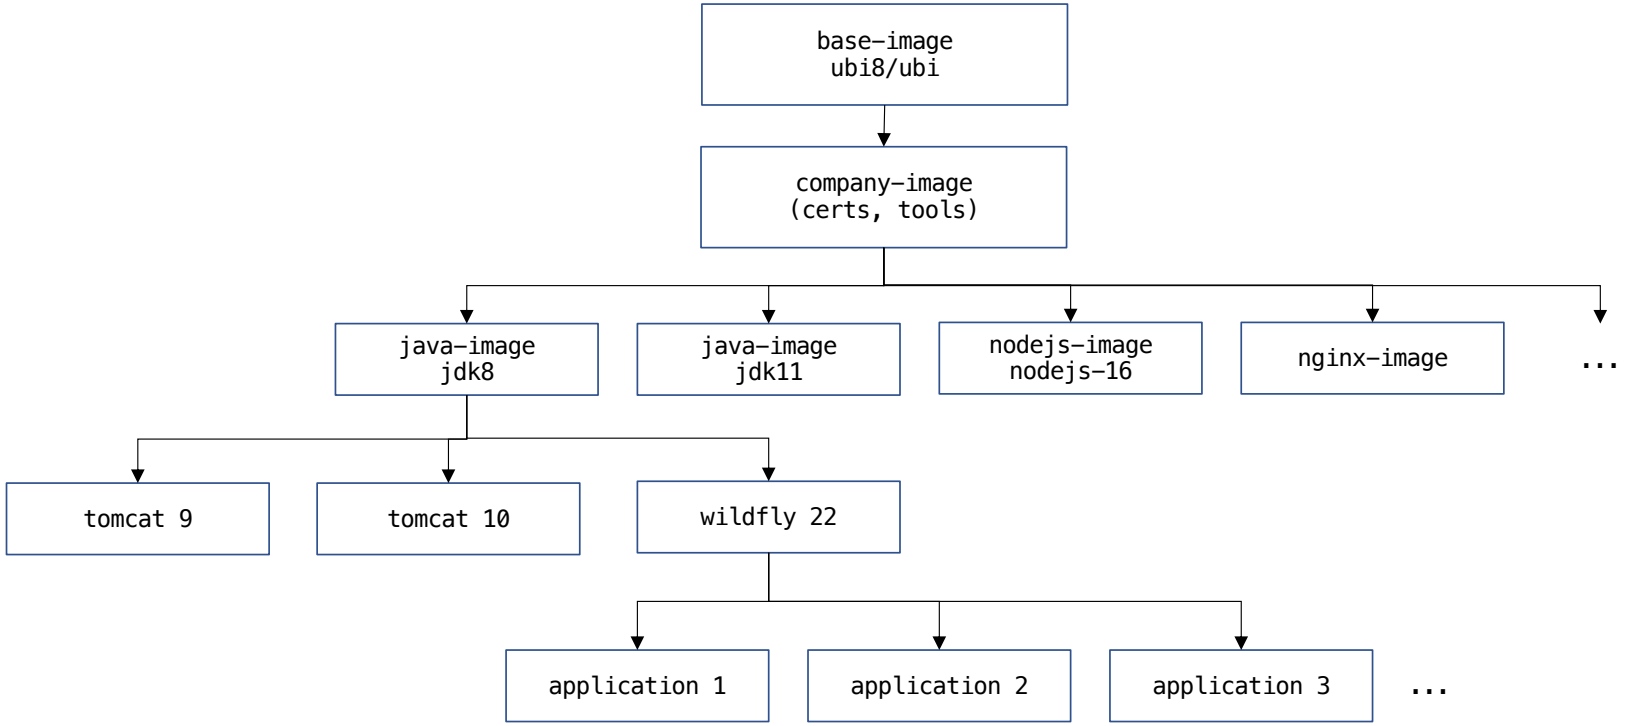


Änderungen an einem Basis-Image erfordern Rebuild der davon abhängigen Images !

image tree

Container in Openshift:

- beliebige User-Id RUN chmod - R 0770 ....
- Group-Id 0 (root) RUN chgrp –R 0
- Ports > 1024

```
apiVersion: project.openshift.io/v1
kind: Project
metadata:
  annotations:
    openshift.io/sa.scc.mcs: s0:c26,c15
    openshift.io/sa.scc.supplemental-groups: 1000680000/10000
    openshift.io/sa.scc.uid-range: 1000680000/10000
```

```
# oc exec pgadmin-778c479f79-tfbqn -- id
uid=1000680000(1000680000) gid=0(root) groups=0(root),1000680000

# ls -al /mnt/nfs/apps/pgadmin
-rw-r--r-- 1 1000680000 root 124K Nov 27 01:03 access_log
-rw-r--r-- 1 1000680000 root  853 Nov 27 00:44 config_local.py
-rw-r--r-- 1 1000680000 root 1.2K Nov 27 00:46 error_log
```

NFS-Mount →

https://cloud.redhat.com/blog/a-guide-to-openshift-and-uids

Abweichende User-Id : Serviceaccount mit Security Context Constraint 'anyuid' notwendig :

```
apiVersion:
rbac.authorization.k8s.io/v1
kind: ClusterRole
metadata:
  name: scc-anyuid
rules:
- apiGroups:
  - security.openshift.io
  resourceNames:
  - anyuid
  resources:
  - securitycontextconstraints
  verbs:
  - use
```

```
apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: gitea:anyuid
  namespace: apps
roleRef:
  kind: ClusterRole
  name: scc-anyuid
  apiGroup: rbac.authorization.k8s.io
subjects:
- kind: ServiceAccount
  name: gitea
  namespace: apps
```

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: gitea
  namespace: apps
```

erstellt von Cluster-Administrator !

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: gitea
  namespace: apps
  ...
spec:
  template:
    spec:
      serviceAccountName: gitea
  ...
```

```
# oc exec gitea-7dcdc5c445-w9qmv -- id
uid=65534(nobody) gid=65534(nobody) groups=65534(nobody),0(root)

# ll /mnt/nfs/repos/ds
drwxr-xr-x 7 nobody nobody 119 Nov 26 16:57 admin.git/
drwxr-xr-x 7 nobody nobody 119 Nov 26 16:12 calibre.git/
drwxr-xr-x 7 nobody nobody 119 Nov 17 16:02 gitea.git/
...
```

UserId aus Container-Config !

## Container Registry:

Red Hat → https://access.redhat.com/RegistryAuthentication

```
# podman login quay.io
Username: ...
Password: ...
Login Succeeded!     -> /run/user/<user-id>/containers/auth.json

# podman push --creds <username>:<password> ...

# skopeo --help
Various operations with container images and container image registries

Usage:
  skopeo [command]

Available Commands:
  copy                                        Copy an IMAGE-NAME from one location to another
  delete                                      Delete image IMAGE-NAME
  help                                        Help about any command
  inspect                                     Inspect image IMAGE-NAME
  list-tags                                   List tags in the transport/repository specified by the REPOSITORY-NAME
  login                                       Login to a container registry
  logout                                      Logout of a container registry
  manifest-digest                             Compute a manifest digest of a file
  standalone-sign                             Create a signature using local files
  standalone-verify                           Verify a signature using local files
  sync                                        Synchronize one or more images from one location to another
```

```
skopeo copy --format ... --dest-creds <user>:<password> containers-storage:localhost/webserver docker://quay.io/do288/webserver
```

## Verwenden einer externen Container Registry - Authentifizierung

```
$ oc get serviceaccounts
NAME       SECRETS
default    2
...
```

```
$ oc create secret docker-registry <secret-name> --docker-server <registry> --docker-username <user> --docker-password <password>

$ oc secrets link <serviceaccount-name> <secret-name> --for pull
```

## Verwenden einer externen Container Registry - Secret von auth.json

```
$ oc create secret generic quayio --from-file .dockerconfigjson=/run/user/1000/containers/auth.json --type kubernetes.io/dockerconfigjson
```

```
apiVersion: v1
kind: Secret
metadata:
  name: quayio
type: kubernetes.io/dockerconfigjson
data:
  .dockerconfigjson: ewogICJhdXRocyI6IHsKICAgICJyZWdpc3 ...
```

### Serviceaccount 'imagePullSecrets' :

```
$ oc secrets link <serviceaccount-name> <secret-name> --for pull
```

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: default
imagePullSecrets:
- name: default-dockercfg-4sdrk
- name: quayio
...
```

oder im Deployment verwenden:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: pgadmin
spec:
  replicas: 1
  template:
    spec:
      imagePullSecrets:
      - name: quayio
      containers:
      - name: pgadmin
        image: registry.connect.redhat.com/crunchydata/crunchy-pgadmin4
```

Imagestream:

- enthält Verweise (Zeiger) auf Images und deren Tags (keine Images)
- Verwendung in Deployment als Image und Trigger
- Import aus externer Registry oder Ergebnis eines Build

```
$ oc import-image webserver --from=quay.io/do288/webserver --confirm  ( --scheduled)
$ oc describe is webserver
Name:                           nginx
Unique Images:                  4
Tags:                           4

latest
  updates automatically from registry quay.io/do288/webserver:latest
  * quay.io/do288/nginx@sha256:c34f57431167fca470730b67a1a8636126d2464eee619ec8d0b577c8e63bffef

1.2
  updates automatically from registry quay.io/do288/webserver:1.2
  * quay.io/do288/nginx@sha256:ee508edacfe0bc1e6af43a15348b400a7d97121507348bd5fb5effb6b9f8d84e

1.1
  updates automatically from registry quay.io/do288/webserver:1.1
  * quay.io/do288/nginx@sha256:674ab485f6e83f162eb4bdaf12986469c7b4f484f65fbb18f3b03218fd5f36e4

1.0
  updates automatically from registry quay.io/
  * quay.io/do288/nginx@sha256:693b30b107da26
```

| TAG | LAST MODIFIED ↓ | SECURITY SCAN | SIZE | MANIFEST |
|---|---|---|---|---|
| 1.2 | 40 minutes ago | 8 Medium | 91.9 MB | SHA256 ee508edacfe0 |
| latest | 14 hours ago | 8 Medium | 91.9 MB | SHA256 c34f57431167 |
| 1.1 | a day ago | 8 Medium | 90.6 MB | SHA256 674ab485f6e8 |
| 1.0 | a day ago | 8 Medium | 90.6 MB | SHA256 693b30b107da |

Verwenden von Imagestreams:

```
$ oc set image-lookup webserver
$ oc describe is webserver | grep Lookup
Image Lookup:                 local=true

$ oc create depoyment webserver --image webserver  ← wird esetzt mit Image-ID

$ oc describe deployments webserver | grep Image:
    Image:  quay.io/do288/webserver@sha256:8c5458214997261f09bd0ffdec0552b57c8c0a8737decf07b51b89997a43c3f4
```

Aktualisierung des Deployments bei Änderungen im ImageStreams:

```
$ oc set triggers deployment webserver --from-image webserver:latest --container webserver

$ oc get deployments.apps webserver -o jsonpath='{.metadata.annotations.image\.openshift\.io/triggers}' | jq
[
  {
    "from": {
      "kind": "ImageStreamTag",
      "name": "webserver:latest"
    },
    "fieldPath": "spec.template.spec.containers[?(@.name==\"webserver\")].image"
  }
]
```

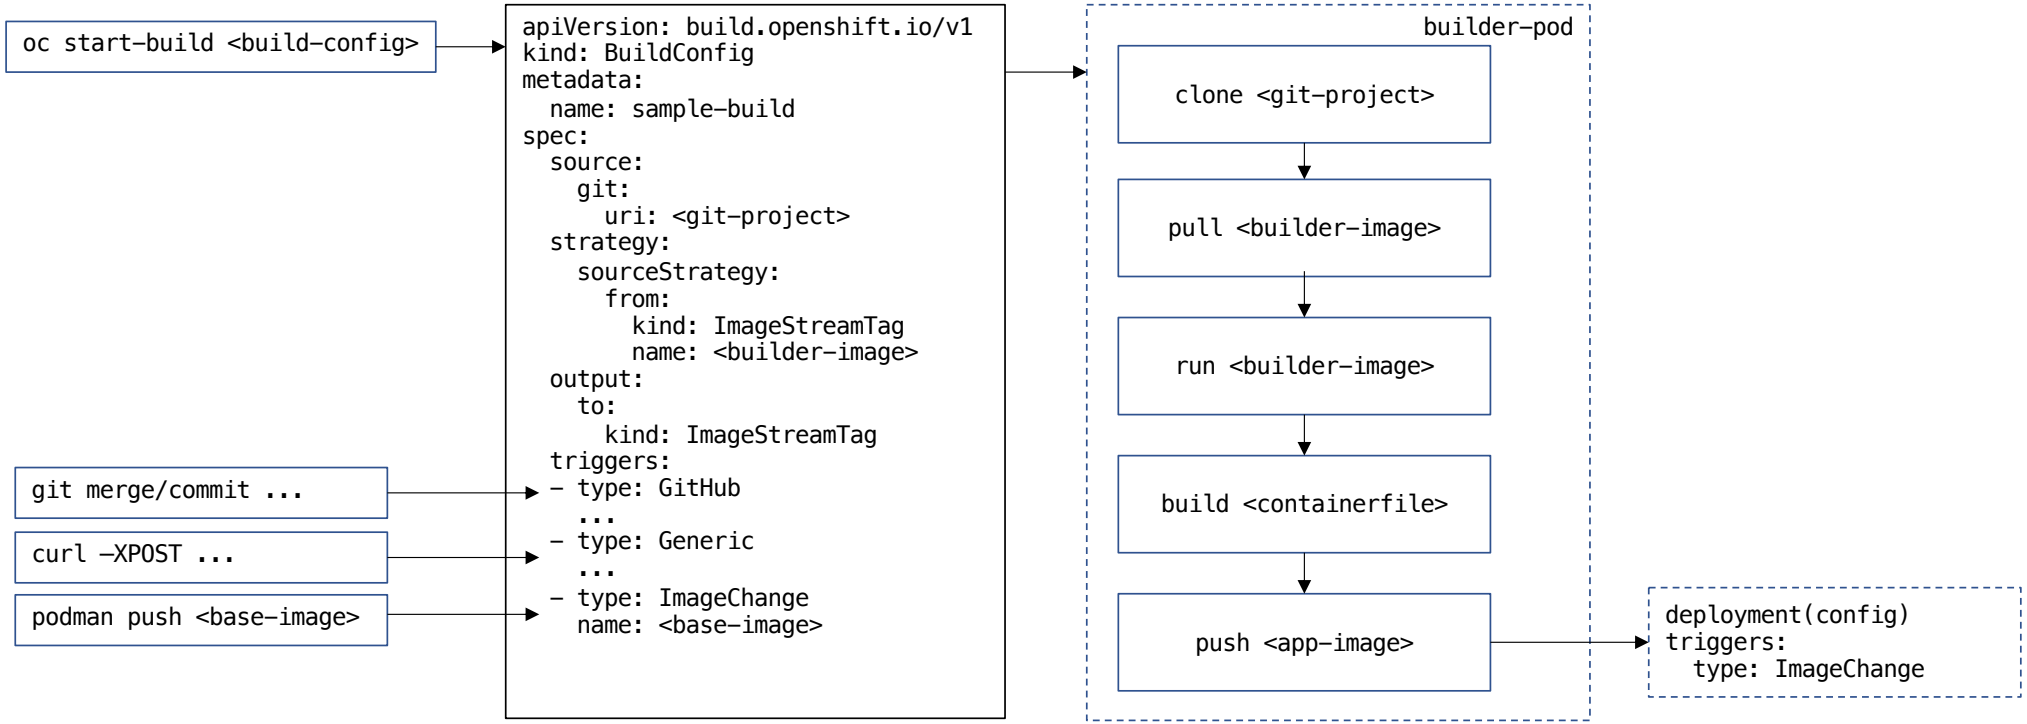


Source: binary | dockerfile | git | images

Strategy:

- source : Builder-Image enthält Tools und Logik zum Erstellen einer Anwendung (Source2Image)
- docker : Git-Repository mit Dockerfile

```
strategy:
  dockerStrategy:
    dockerfilePath: Containerfile
    ....
```

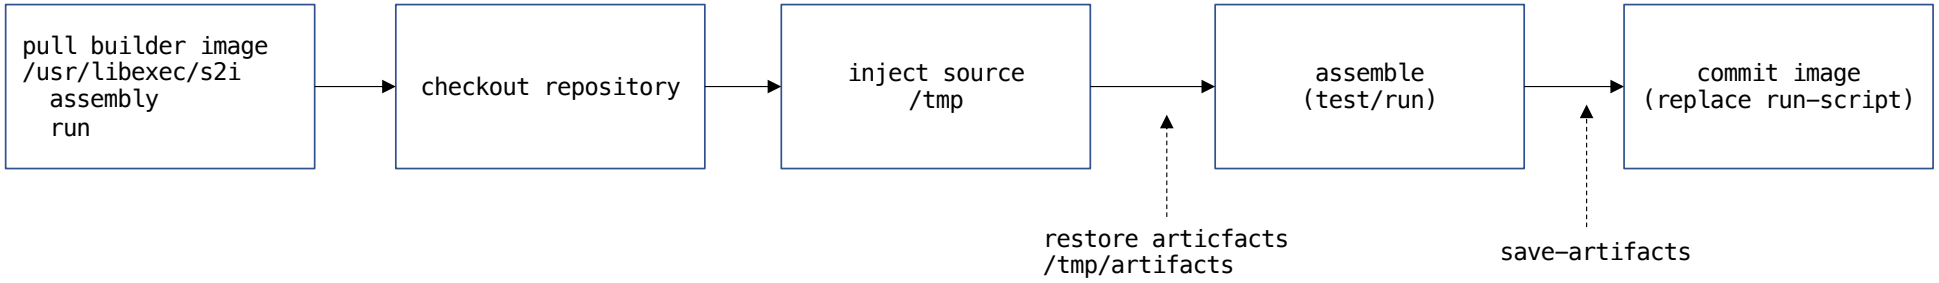


Build-Scripte:

- default in `/usr/libexec/s2i`
- `assemble` und `run` sind mandatory
- `save-artifacts`, usage, `test/run` sind optional
- könnnen überschrieben werden im Git-Repo `.s2i/bin`
  (Wrapper um Original-Script oder komplett neues Script)

Incremental Builds:

- `save-artifacts` erstellt TAR
- wird vor dem Ausführen von assembly injected in `/tmp/artifacts`

Deployment-Strategien

- Rolling Updates : Pods werden der Reihe nach aktualisiert
- Recreate: existierende Pods werden beendet und neue gestartet

DeploymentConfig: DEPRECATED ab 4.15 !

- Pre/Mid/Post – Lifecycle Hooks

Beenden eines Pods:

- SIGTERM: Pod soll keine neuen Verbindungen annehmen und bestehenden Aktionen beenden
- SIGKILL: nach `terminationGracePeriodSeconds` (30s) wird der Pod beendet

```
kind: Deployment
metadata:
  name: ...
spec:
  revisionHistoryLimit: 3   (default: 10)
  replicas: 4
  strategy:
    type: RollingUpdate
    rollingUpdate:
      maxSurge: 1             ← max. 5 Pods aktiv
      maxUnavailable: 1
  ...
  template:
    spec:
      containers:
      - ...
      terminationGracePeriodSeconds: 30
```

```
oc rollout SUBCOMMAND deployment <name>

  cancel      Cancel the in-progress deployment
  history     View rollout history
  latest      Start a new rollout for deployment config with latest state
  pause       Mark the provided resource as paused
  restart     Restart a resource
  resume      Resume a paused resource
  retry       Retry the latest failed rollout
  status      Show the status of the rollout
  undo        Undo a previous rollout
```

N-1 Abwärtskompatibilität bei Rolling-Update:

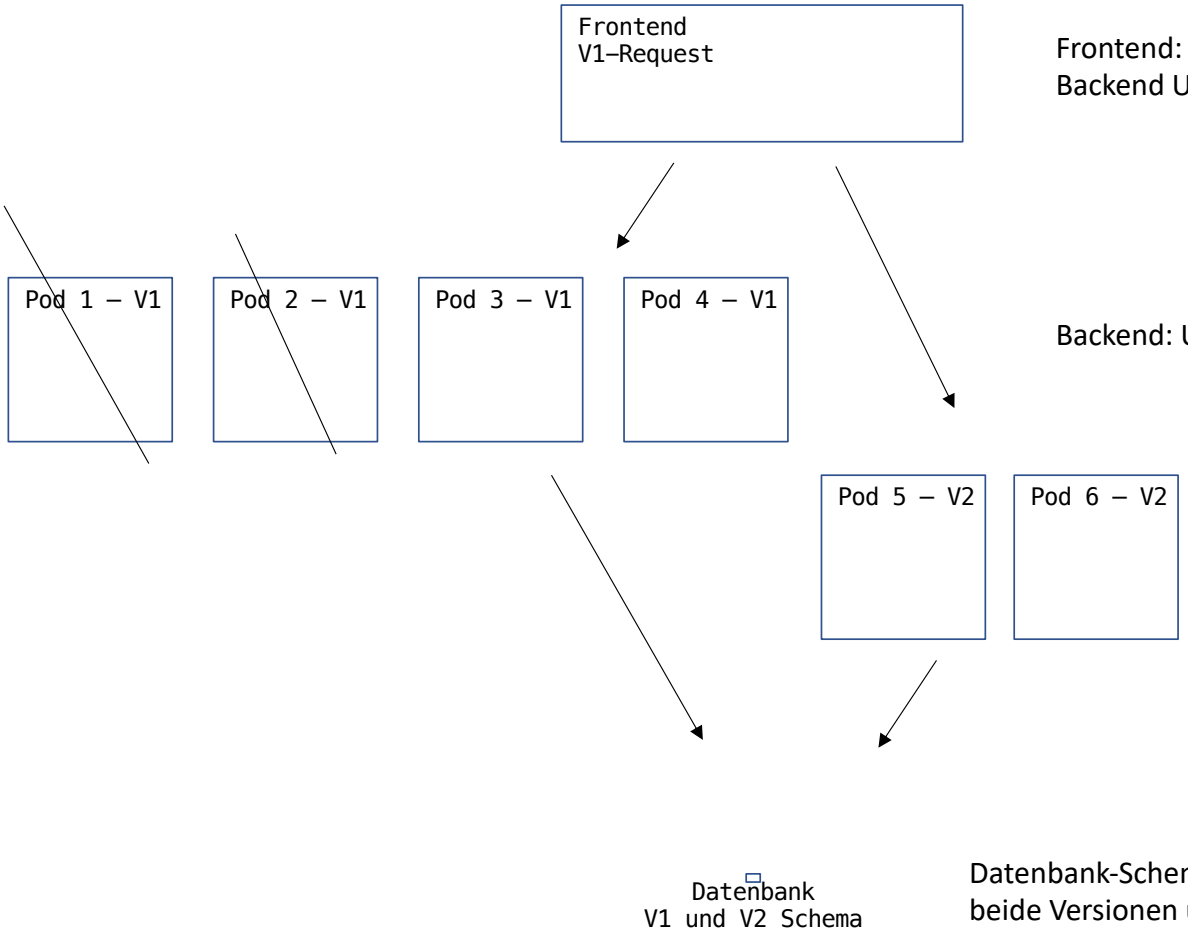


Frontend: behält V1 bis Backend Update fertig ist

Backend: Update von V1 auf V2

Datenbank-Schema muss auch beide Versionen unterstützen

## A/B Deployment Strategy:

```yaml
apiVersion: v1
kind: Service
metadata:
  name: service-a
spec:
ports:
  - name: http
    port: 80
    protocol: TCP
    targetPort: http
  selector:
    app.kubernetes.io/instance: deploment-a
```

```yaml
apiVersion: v1
kind: Service
metadata:
  name: service-b
spec:
ports:
  - name: http
    port: 80
    protocol: TCP
    targetPort: http
  selector:
    app.kubernetes.io/instance: deploment-b
```

```yaml
kind: Route
metadata:
  name: <name>
spec:
  host: <host>
  to:
    kind: Service
    name: service-a
    weight: 198
  alternateBackends:
  - kind: Service
    name: service-b
    weight: 2
```

## Secrets:

- Passwörter, Token, Zertifikate ...
- typisiert: basic-auth, dockerfg, tls, opaque
- Inhalte sind base64-decodiert, nicht verschlüsselt

## ConfigMap:

- generische Key-Value Daten

→ max. Größe 1 MB
→ nur innerhalb eines Project (NS) sichtbar

```
apiVersion: v1
kind: Secret
metadata:
  name: ...
  namespace: ...
data:
  password: MTIzNDU2
type: Opaque
```

```
# echo MTIzNDU2 | base64 -d
123456
```

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: ...
  namespace: ...
binaryData:
  keystore: |
    7oAMCAQICCF7Dt6ZDf6TgMA0GCSqGSIb3DQEBBQUAMEI1ZSQUla
    MTEQMA4GA1UECwwHU ...
data:
  HOME: /usr/share/nginx
  default.conf: |
    server {
      listen 8181 default_server;
      server_name _;
      location / {
         root   /usr/share/nginx/html;
         index  index.html index.htm;
      }
    }
```

```
$ oc create configmap <cm-name> --from-literal FOO=BAR

$ oc create configmap <cm-name> --from-file <path>

$ oc create secret docker-registry quayio --docker-server quay.io --docker-username <user> --docker-password <password>
```

## Secrets: Verwendung als Umgebungs-Variable

```
apiVersion: v1
kind: Pod
metadata:
  name: secret-env-pod
spec:
  containers:
  - name: mycontainer
    image: redis
    env:
    - name: SECRET_USERNAME
      valueFrom:
        secretKeyRef:
           name: mysecret
           key: username
    - name: SECRET_PASSWORD
      valueFrom:
        secretKeyRef:
           name: mysecret
           key: password
```

## ConfigMap: Verwendung als Konfigurations-Dateien

```
apiVersion: apps/v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    container: nginx
    volumeMounts:
    - mountPath: /etc/nginx/conf.d
      name: config
  volumes:
   - name: config
     configMap:
       name: nginx-config
```

```
apiVersion: apps/v1
kind: Pod
metadata:
  name: wildfly-standalone-xml
spec:
  containers:
  - name: wildfly
    container: nginx
    volumeMounts:
    - mountPath: /opt/wildfly/standalone/configuation
      name: standalone-xml
      subPath: standalone.xml
  volumes:
   - name: standalone-xml
     configMap:
       name: standalone-xml
```

```
$ oc set env deployment/<deployment-name> --from cm/<cm-name>

$ oc set volume deployment/<deployment-name> -add -t configmap -m /etc/nginx/conf.d --name config --configmap-name <cm-name>
```

## Sealed Secrets –> https://github.com/bitnami-labs/sealed-secrets

```
$ oc create secret generic sample --from-literal user=daniel ... --dry-run=client -o yaml | kubeseal -o yaml
```

```
apiVersion: bitnami.com/v1alpha1
kind: SealedSecret
metadata:
name: sample
  namespace: <namespace>
spec:
  encryptedData:
    password: AgBkrNLR....
    user: AgCvHr....
  template:
    metadata:
      name: sample
      namespace: <namespace>
```

```
$ oc apply -f sample.yml
$ oc get secrets

NAME            TYPE                DATA
sample          Opaque              2
...

$coc get sealedsecrets.bitnami.com
NAME     STATUS   SYNCED
sample            True
```

**External Secrets → https://external-secrets.io**

| | |
|---|---|
| (Cluster) SecretStore | Verbindung zu einer "Secret-Quelle" (Vault, Namespace) |
| (Cluster) ExternalSecret | Erzeugt Secret aus einer Referenz vom SecretStore |

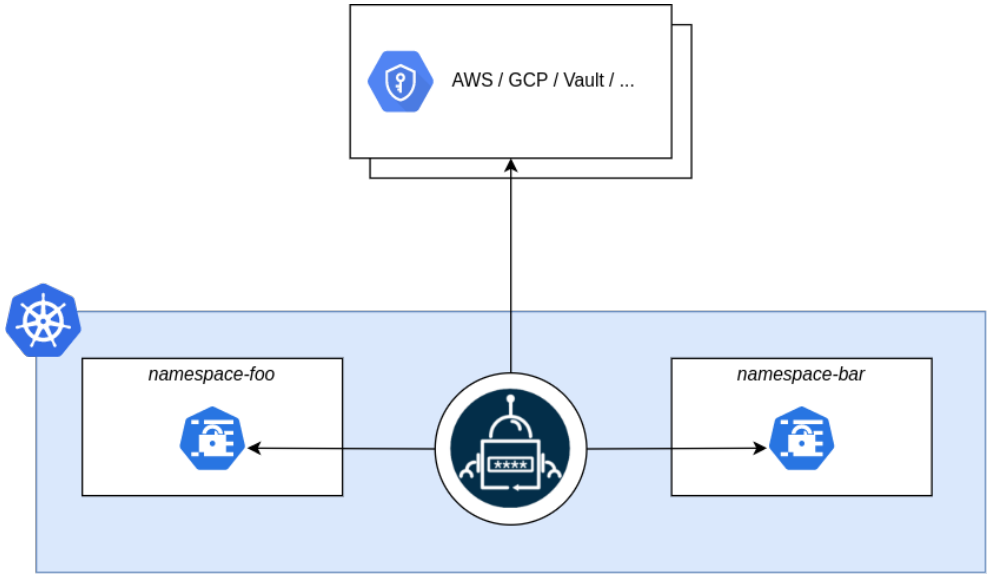

## **External Secrets -** Beispiel Namespace als Secret-Quelle

```yaml
apiVersion: external-secrets.io/v1beta1
kind: ClusterSecretStore
metadata:
   name: external-secrets
spec:
  provider:
    kubernetes:
      auth:
        serviceAccount:
          name: secrets-reader
          namespace: external-secrets
      remoteNamespace: external-secrets
      server:
        url: kubernetes.default
        ...
```

```yaml
apiVersion: external-secrets.io/v1beta1
kind: ClusterExternalSecret
metadata:
  name: registry-conf
spec:
  refreshInterval: 1m
  namespaceSelector:
    matchLabels:
      ext-secret-registry-conf: "true"
  externalSecretSpec:
    secretStoreRef:
     ....
```

ns: external-secrets

secret git-ssh

secret ssl-certs

secret registry-conf

ns: webserver

```yaml
apiVersion: v1
kind: Namespace
metadata:
 name: webserver
 labels:
    ext-secret-registry-conf: "true"
```

```yaml
apiVersion: v1
kind: Secret
metadata:
  annotations:
    reconcile.external-secrets.io/data-hash: 8bb9...
labels:
    reconcile.external-secrets.io/created-by: f3a4...
  name: registry-conf
  namespace: webserver
  ownerReferences:
  - apiVersion: external-secrets.io/v1beta1
    kind: ExternalSecret
data:
  ...
```

# Volumes

```
node

/
├── bin
├── etc
├── lib        mount
├── mnt    ──────────►  container
├── opt                 /
├── proc                ├── bin
├── usr                 ├── etc
└── var                 ├── lib
                        ├── mnt
                        ├── opt
                        ├── proc
                        ├── usr
                        └── var

podman –v <local>:<container>
```

```
kind: Pod
...
spec:
  containers:
     ...
     volumeMounts:
       – mountPath: <path_container_fs>
         name: <name>
      ...
  volumes:
     – name: <volume>
       <volume–type>:
          <volume–attributes>
```

→ https://kubernetes.io/docs/concepts/storage/volumes/

Volume=Types
- emptyDir
- hostPath (system:openshift:scc:hostmount–anyuid !)
- configMap
- secret
- peristentVolumeClaim

...

# Persistence

## Administrator erzeugt PersistentVolume

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: nfs-data
spec:
  accessModes:
    - ReadWriteMany
  capacity:
    storage: 10Gi
  nfs:
    path: /mnt/nfs/data
    server: 10.0.0.1
  persistentVolumeReclaimPolicy: Retain
```

- automatisiertes PV-Management mit storageClass/Provisioner

## Anwendung erstellt Anforderung

```
kind: PersistentVolumeClaim
apiVersion: v1
metadata:
  name: html-data
spec:
  accessModes:
    - ReadWriteMany
  resources:
    requests:
      storage: 10Gi
```

## und verwendet dieses im Deployment / Pod

```
kind: Deployment
...
    containers:
      - name: webserver
        ...
        volumeMounts:
         - mountPath: /usr/share/ngnix/html
           name: html
    volumes:
      - name: html
        persistentVolumeClaim:
          claimName: html-data
```

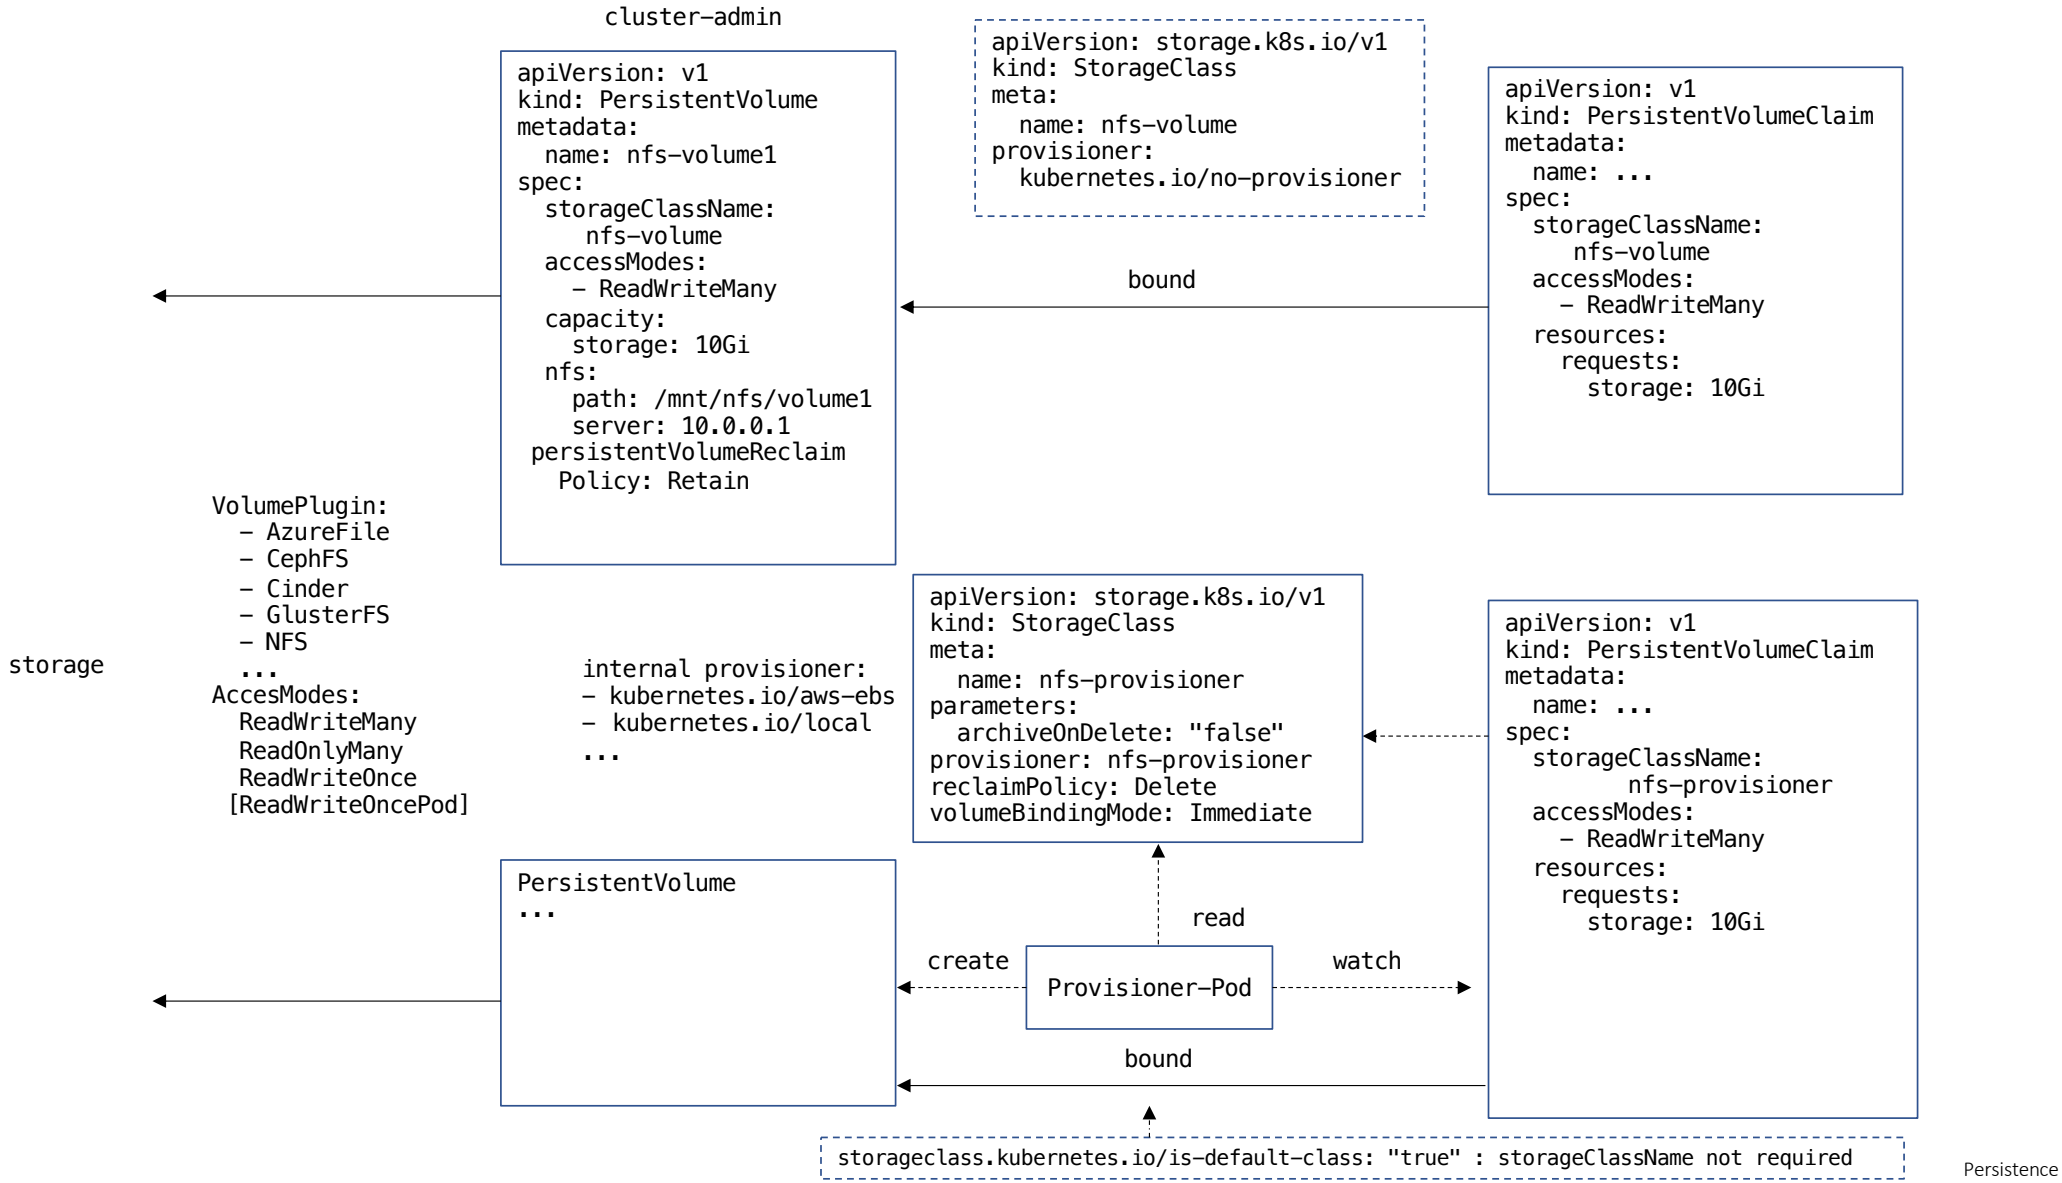


Persistence

# StatefulSet

```yaml
apiVersion: v1
kind: Service
metadata:
  name: <svc-name>
spec:
  clusterIP: None
  ports:
  - name:
    ...
  selector:
    app: <name>
---
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: <name>
spec:
  serviceName: <svc-name>
  replicas: 3
  ...
  template:
    metadata:
      labels:
        app: <name>
    spec:
      containers:
      - name: ...
        ...
        volumeMounts:
        - name: <pvc-name>
          mountPath: ...
  volumeClaimTemplates:
  - metadata:
      name: <pvc-name>
    spec:
      accessModes:
      - ReadWriteMany
      resources:
        requests:
          storage: xxGi
```

***Headless-Service*** (keine Cluster-IP, kein Loadbalancing)

Host-Name/DNS A-Record für jeden Pod: <name-#>.<svc-name>.<project>.svc.cluster.local

CNAME <svc-name>.<project>.svc.cluster.local + SVR-Records für jeden Pod

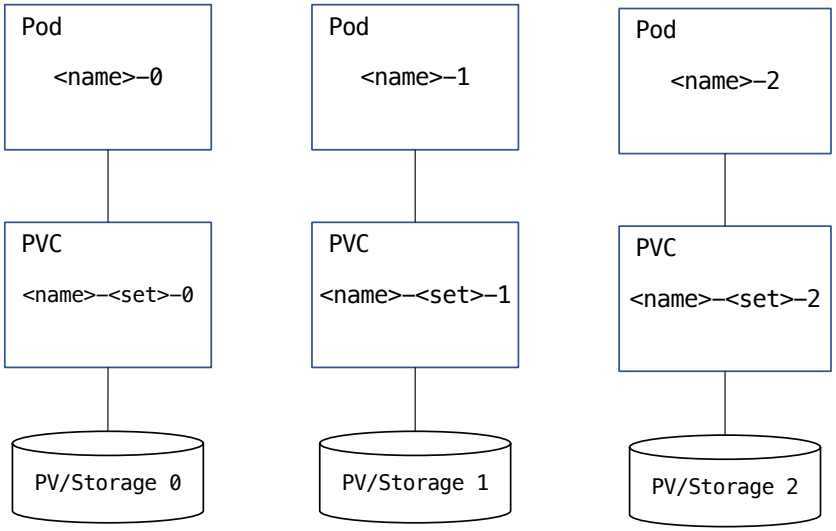

(compute) Resources :

- Memory: number of bytes (quantity suffixes: E, P, T, G, M, k | Ei, Pi, Ti, Gi, Mi, Ki)
- CPU : millicores (m) … fractions of *time* of a single CPU (not the fraction of number of CPUs).

```
apiVersion: v1
kind: Pod
metadata:
...
spec:
  containers:
  - name: <name>
    resources:
      requests:
        memory: 64Mi
        cpu: 100m
      limits:
        memory: 128Mi
        cpu: 200m
```

Scheduling

Execution
(cgroups)

```
$ oc describe node master01
...
Allocatable:
  cpu:                3500m
  memory:             15268156Ki
Non-terminated Pods: (60 in total)
   CPU Requests CPU Limits Memory Requests Memory Limits
...
Allocated resources:
Resource           Requests      Limits
  --------           --------      ------
  cpu                2397m (68%)   0 (0%)
  memory             9347Mi (62%)  512Mi (3%)
```

resouce usage

```
$ oc get pods -w
NAME                                      READY   STATUS              RESTARTS   AGE
famousapp-famouschart-65744d4c8b-4zqhn    0/1     Running             0          10s
famousapp-mariadb-0                       0/1     ContainerCreating   0          10s
famousapp-mariadb-0                       0/1     Running             0          11s
famousapp-famouschart-65744d4c8b-4zqhn    0/1     Running             1          33s
famousapp-famouschart-65744d4c8b-4zqhn    0/1     Error               1          34s
famousapp-famouschart-65744d4c8b-4zqhn    0/1     Running             2          35s
famousapp-famouschart-65744d4c8b-4zqhn    0/1     Error               2          36s
famousapp-famouschart-65744d4c8b-4zqhn    0/1     CrashLoopBackOff    2          37s
famousapp-mariadb-0                       1/1     Running             0          48s
famousapp-famouschart-65744d4c8b-4zqhn    0/1     Running             3          56s
famousapp-famouschart-65744d4c8b-4zqhn    1/1     Running             3          62s
```

```
metadata:
  name: famousapp-mariadb:
...

  livenessProbe:
    exec:
      command:
        - /bin/bash
        - -ec
        - |
          password_aux="${MARIADB_ROOT_PASSWORD:-}"
          if [[ -f "${MARIADB_ROOT_PASSWORD_FILE:-}"
             password_aux=$(cat "$MARIADB_ROOT_PASSW
          fi
          mysqladmin status -uroot -p"${password_aux
```

```
metadata:
  name: famousapp-famouschart
...
   livenessProbe:
      initialDelaySeconds: 30
      httpGet:
        path: /
      ...
   readinessProbe:
      failureThreshold: 3
      httpGet:
        path: /
        port: http
        scheme: HTTP
        periodSeconds: 10
        successThreshold: 1
        timeoutSeconds: 1
```

## Liveness / Readiness / Startup Probes

- liveness : Container wird bei negativen Ergebnis neu gestartet `.spec.containers.livenessProbe`
- readiness: Route/Service wird aktiviert/deaktiviert `.spec.containers.readinessProbe`
- startup: liveness/readiness sind deaktiviert bis startup positiv ist `.spec.containers.startupProbe`
  Container wird bei neg. Startup-Probe sofort beendet

## Probes:

```
exec:
  command:
  - cat
  - /tmp/ready
initialDelaySeconds: 5
periodSeconds: 5
```

```
httpGet:
  path: /healthz
  port: healthz-port
  schema: https
  httpHeaders: ...
failureThreshold: 1
periodSeconds: 10

200 <= status < 400
```

```
tcpSocket:
  port: 5432
initialDelaySeconds: 15
periodSeconds: 20
```

- initialDelaySeconds: Zeitdauer bis zur ersten liviness/readiness Probe
- periodSeconds: Intervall zur Ausführung der Proben (default 10 sec)
- timeoutSeconds: max. Timeout bei einer Probe (default 1 sec)
- successThreshold: Schwellwert ab wann aufeinderfolgende positive Proben als Erfolg gewertet werden (default 1)
- failureThreshold: Schwellwert ab wann aufeinderfolgende negative Proben als Ausfall gewertet werden (default 3)

```yaml
kind: Deployment
apiVersion: apps/v1
metadata:
  name: webserver
spec:
  ...
  template:
    spec:
      containers:
      - name: webserver
        image: webserver
        imagePullPolicy: Always
        ports:
        - name: http
          containerPort: 8080
          protocol: TCP
        readinessProbe:
          failureThreshold: 3
          httpGet:
            path: /healthz
            port: http
            scheme: HTTP
          periodSeconds: 10
          successThreshold: 1
          timeoutSeconds: 1
        ...
```

nginx.conf

```
server {
  listen 8080 default_server;
  server_name _;
  location / {
    root   /usr/share/nginx/html;
    index  index.html index.htm;
  }

  location /healthz {
    access_log off;
    return 200;
  }

}
```

# Pod – Scheduling

1. Filter
   - Node-Selector für Labels
     https://kubernetes.io/docs/reference/labels-annotations-taints

   - Toleration für Taints
     https://kubernetes.io/docs/concepts/scheduling-eviction/taint-and-toleration

```
apiVersion: v1
kind: Pod
metadata:
...
spec:
  containers:
  - name: nginx
    nodeSelector:
      disktype: ssd
  tolerations:
  - key: class
    value: do280
    operator: "Equal"
    effect: "NoSchedule"
```

```
apiVersion: v1
kind: Node
metadata:
 labels:
   disktype: ssd
spec:
  taints:
  - key: class
    value: do280
    effect: NoSchedule
```

Pod – Scheduling

2. Scoring
   Affinity/Anti-Affinity-Rules

```
apiVersion: v1
kind: Pod
metadata:
  name: with-node-affinity
spec:
  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        nodeSelectorTerms:
        - matchExpressions:
          - key: kubernetes.io/os
            operator: In
            values:
            - linux
      preferredDuringSchedulingIgnoredDuringExecution:
      - weight: 1
        preference:
          matchExpressions:
          - key: another-node-label-key
            operator: In
            values:
            - another-node-label-value
  containers:
  - name: with-node-affinity
    image: ...
```

...DuringScheduling: während des Scheduling

IgnoredDuringExecution: Pod wird weiter ausgeführt,
auch wenn sich nach dem Scheduling Node-Labels ändern

https://kubernetes.io/docs/concepts/scheduling-eviction/assign-pod-node
https://www.cncf.io/blog/2021/07/27/advanced-kubernetes-pod-to-node-scheduling

Pod Verteilung auf unterschiedliche Nodes:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  ...
spec:
  selector:
    matchLabels:
      app: store
  replicas: 3
  template:
    metadata:
      labels:
        app: store
    spec:
      affinity:
        podAntiAffinity:
          requiredDuringSchedulingIgnoredDuringExecution:
          - labelSelector:
              matchExpressions:
              - key: app
                operator: In
                values:
                - store
            topologyKey: "kubernetes.io/hostname"
      containers:
      - ...
```

gleiche Label und gleicher Hostname → AntiAffinity

Pods werden auf unterschiedlichen Nodes verteilt

https://docs.openshift.com/container-platform/4.10/nodes/scheduling/nodes-scheduler-pod-affinity.html#nodes-scheduler-pod-affinity-example-antiaffinity_nodes-scheduler-pod-affinity


Multi Container App / Staging

Template: parametrisierbare Liste von Resource-Definitionen

```yaml
kind: Template
apiVersion: v1
metadata:
  name: rest-sample
objects:
- apiVersion: v1
  kind: Service
  metadata:
    name: ${APP_NAME}
  spec:
    selector:
      app.kubernetes.io/name: ${APP_NAME}
   ...
- apiVersion: apps/v1
  kind: Deployment
  metadata:
    name: ${APP_NAME}
  spec:
    template:
      spec:
        containers:
        - name: ${APP_NAME}
          image: ${IMAGE_NAME}
...
- apiVersion: v1
  kind: Route
...
parameters:
- description: Application Name
  name: APP_NAME
  required: true
- description: Image Name
  name: IMAGE_NAME
  required: true
...
```

```
oc process (TEMPLATE | -f FILENAME) -p APP_NAME=... | oc create -f -

oder bei installiertem Template ( oc create -f template.yml) :

oc new-app <template-name>
```

```
apiVersion: template.openshift.io/v1
kind: Template
metadata:
  name: webserver
  labels:
    template: webserver
labels:
  app: webserver
parameters:
- name: stage
  required: true
- name: version
  value: latest
objects:
- apiVersion: v1
  kind: Deployment
  metadata:
    name: webserver
    namespace: stage-${stage}
    ...
       containers:
        - name: webserver
          image: registry.straubcloud.de/webserver:${version}
...
- apiVersion: v1
  kind: Route
  ...
  spec:
    host: webserver-${stage}.ocp.straubcloud.de
  ...
```

- **Namespace stage-test**

```
$ oc process -f webserver.yml -p stage=test | oc apply -f -
deployment.apps/webserver created
service/webserver created
route.route.openshift.io/webserver created
```

- **Namespace stage-prod**

```
$ oc process -f webserver.yml -p stage=prod -p version=2.0 | oc apply -f -
deployment.apps/webserver created
service/webserver created
route.route.openshift.io/webserver created
```

Helm-Chart:  Paket-Manager (Lifecycle + Template-Engine + Dependencies)

```
$ helm create sample
Creating sample

$ tree sample
sample
├── charts
├── Chart.yaml
├── templates
│   ├── deployment.yaml
│   ├── _helpers.tpl
│   ├── hpa.yaml
│   ├── ingress.yaml
│   ├── NOTES.txt
│   ├── serviceaccount.yaml
│   ├── service.yaml
│   └── tests
│       └── test-connection.yaml
└── values.yaml
```

helm create

Helm-Chart: Paket-Manager (Lifecycle + Template-Engine + Dependencies)

Chart.yml

```
apiVersion: v1
name: sample
description: Sample Application
version: 1.0
appVersion: 1.0
dependencies:
- name: dep1
  version: ...
  repository: ...
```

values.yml

```
image:
  repository: quay.io/redhat.io/sample
  tag: '2'

service:
  port: 8080

env:
  ...


dep1.key: value
```

Templates:

deployment.yml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: {{ APP_NAME }}
spec:
  template:
    selector:
      matchLabels:
        {{- include "sample.selectorLabels" . | nindent 6 }}
     spec:
       containers:
         - image: {{.Values.image.repository}}: {{.Values.image.tag}}
 ...
```

Go-Templates:

_helpers.tpl

```
{{- define "sample.selectorLabels" -}}
app.kubernetes.io/name: {{ include "sample.name" . }}
app.kubernetes.io/instance: {{ .Release.Name }}
{{- end }}
...
```

```
helm create
helm dependency update
helm install / upgrade / rollback / uninstall

helm template (lokales processing)
```

Helm: Verwendung unterschiedlicher Stages

```
webserver
├── Chart.yaml
├── stage-prod
│   ├── env.properties
│   └── values.yml
├── stage-test
│   ├── env.properties
│   └── values.yml
├── templates
  ...
└── values.yaml
```

➢ **Namespace stage-test**

```
$ helm install webserver . -f stage-test/values.yml -n stage-test
NAME: webserver
LAST DEPLOYED: ...
NAMESPACE: stage-test
STATUS: deployed
REVISION: 1

$ helm upgrade webserver . -f stage-test/values.yml --set image.tag=1.0
Release "webserver" has been upgraded. Happy Helming!
NAME: webserver
LAST DEPLOYED: ...
NAMESPACE: stage-test
STATUS: deployed
REVISION: 2
```

➢ **Namespace stage-prod**

```
$ helm install webserver . -f stage-prod/values.yml -n stage-prod
NAME: webserver
LAST DEPLOYED: ....
NAMESPACE: stage-prod
STATUS: deployed
REVISION: 1
```

Kustomize: generieren/transformieren von Resourcen (Manifeste mit minimalen Meta-Daten)

kustomization.yml

```
kind: Kustomization
apiVersion: kustomize.config.k8s.io/v1beta1

namespace: sample

resources:
- deployment.yml
- service.yml
- route.yml
- http://...    -> kustomize.yml in Git/Web-Repository


images:
- name: sample
  newName: registry/sample
  newTag: '5'

commonLabels:
  app.kubernetes.io/instance: sample

configMapGenerator:
- name: rest-sample
  literals:
  - LAUNCH_JBOSS_IN_BACKGROUND=1
...
```

resources → https://github.com/hashicorp/go-getter#url-format

deployment.yml

```
apiVersion: apps/v1
metadata:
  name: rest-sample
spec:
  replicas: 1
  template:
    spec:
      containers:
      - name: sample
        image: sample
```

```
$ oc kustomize <kustom-dir>
apiVersion: apps/v1
kind: Deployment
metadata:
  labels:
    app.kubernetes.io/instance: rest-sample
  name: rest-sample
  namespace: sample
spec:
  replicas: 1
  selector:
    matchLabels:
      app.kubernetes.io/instance: sample
  template:
    containers:
      image: registry/sample:5
...

$ oc apply -k .
```

Kustomize Overlays : erzeugen unterschiedlicher Varianten von einer Basis-Vorlage

```
                                        base/kustomization.yml
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization

resources:
- deployment.yml
- service.yml
- route.yml
```

```
                                overlays/test/kustomization.yml
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization

resources:
- ../../base

namespace: test

images:
- name: sample
  newName: registry/sample
  newTag: '3-SNAPSHOT'
```

```
                          overlays/production/kustomization.yml
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization

resources:
- ../../base

namespace: production

images:
- name: sample
  newName: registry/sample
  newTag: '5'
```

```
$ oc apply -k overlays/test
service/sample configured
deployment.apps/sample configured
route.route.openshift.io/sample configured

$ oc apply -k overlays/production
...
```

```
$ grep -A3 images kustomization.yml
images:
- name: webserver
  newName: quay.io/danielstraub/webserver
  newTag: "1.0"

$ kustomize edit set image webserver=quay.io/danielstraub/webserver:2.0

$ grep -A3 images kustomization.yml
images:
- name: webserver
  newName: quay.io/danielstraub/webserver
  newTag: "2.0"

$ oc apply -k .
configmap/webserver-kt5mdg45d2 unchanged
service/webserver unchanged
deployment.apps/webserver configured
route.route.openshift.io/webserver unchanged

$ curl https://stage-prod.apps.eu46a.prod.ole.redhat.com
Hello, DO288
Version 2.0
```

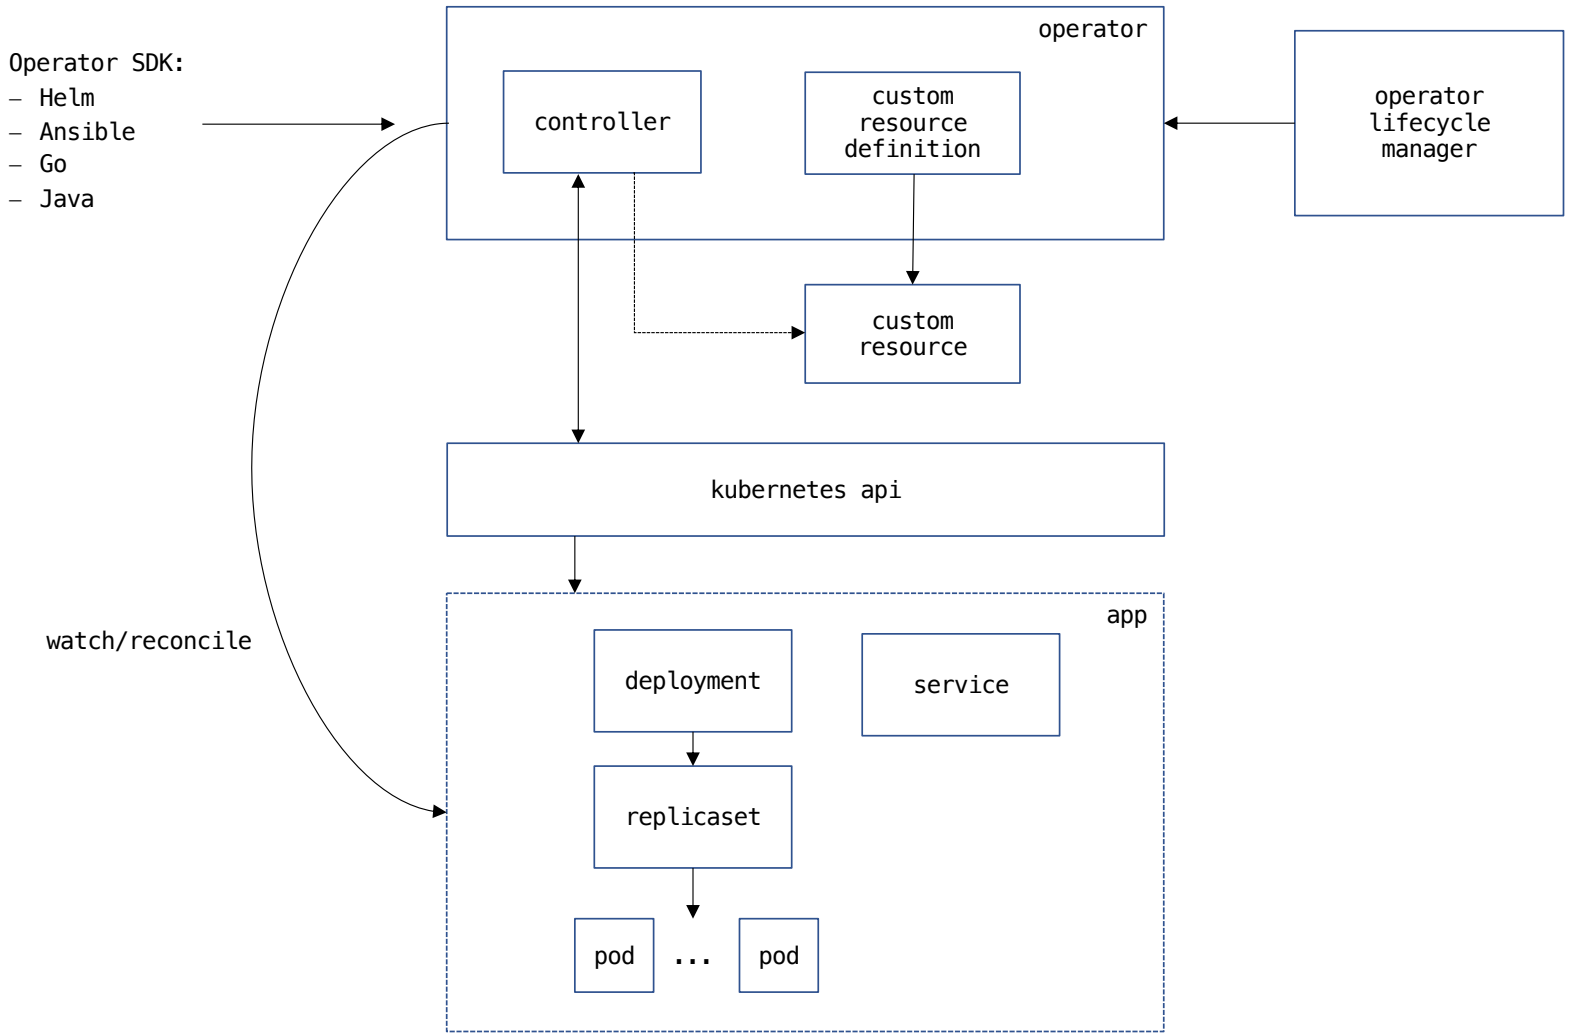
Operator SDK:
– Helm
– Ansible
– Go
– Java
operator
controller
custom resource definition
operator lifecycle manager
custom resource
kubernetes api
watch/reconcile
app
deployment
service
replicaset
pod
...
pod

operator

```
$ oc api-resources | grep keycloak
```

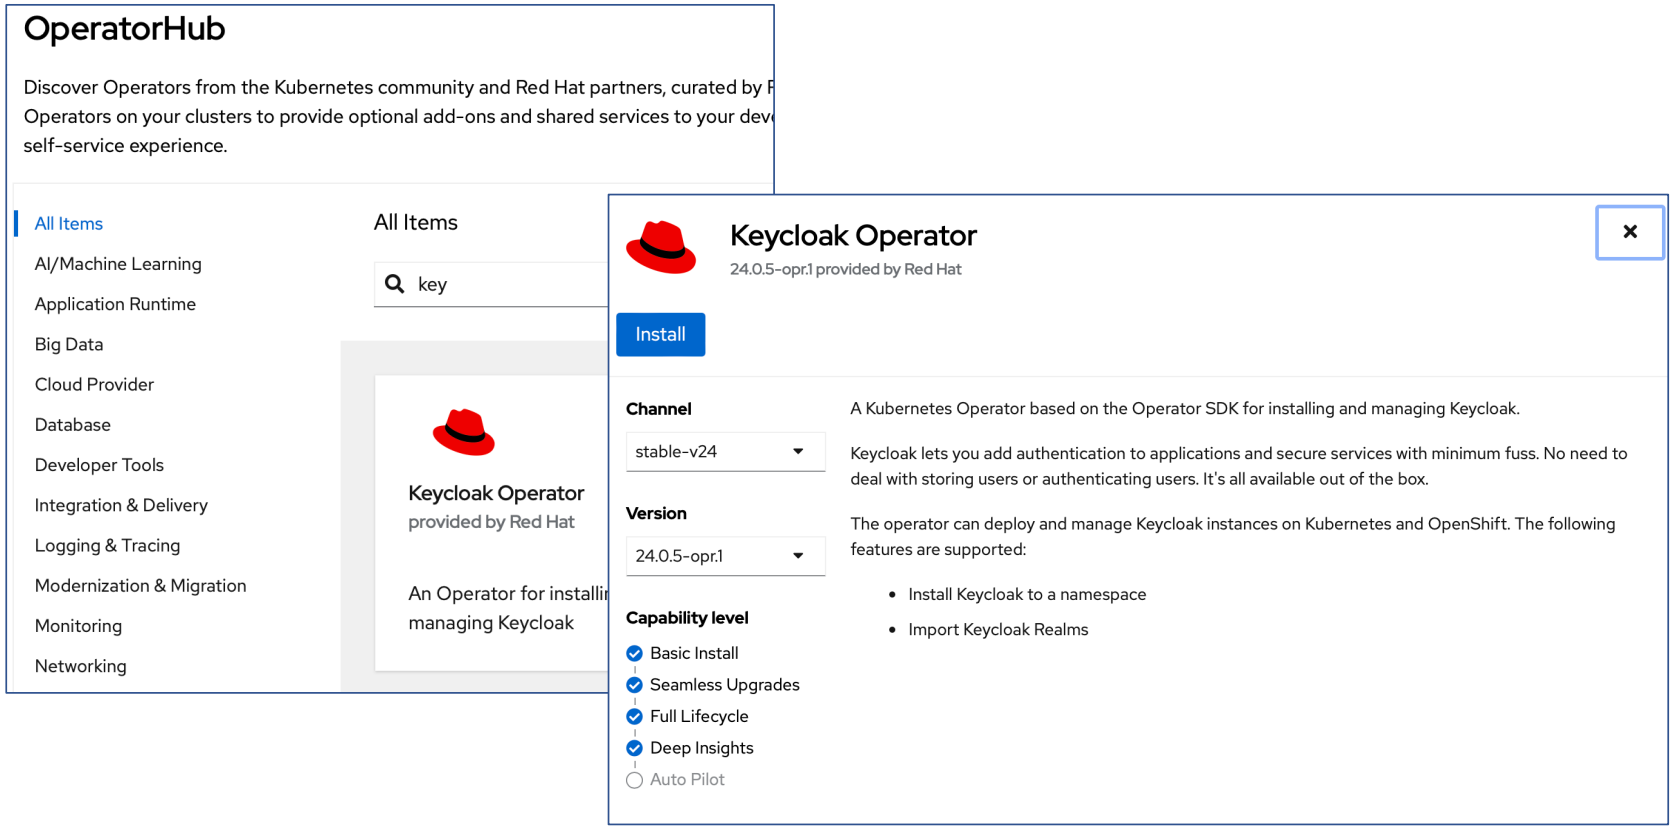


```
oc api-resources | grep keycloak
keycloakrealmimports
keycloaks
```

operator keycloak

```
$ oc project keycloak
$ oc create secret tls keycloak-cert --cert=...fullchain.pem --key=..privkey
```

```
$ oc apply -f - <<EOF
kind: Keycloak
apiVersion: k8s.keycloak.org/v2alpha1
metadata:
  name: keycloak
  labels:
    app: keycloak
  namespace: keycloak
spec:
  instances: 1
  hostname:
    hostname: idp.<wildcard-domain>
  http:
    tlsSecret: keycloak-cert
EOF
```

```
$ oc get all
NAME                                  READY   STATUS    RESTARTS   AGE
pod/keycloak-0                        1/1     Running   0          5m28s
pod/rhbk-operator-565b768dcd-qgqss    1/1     Running   0          13m

NAME                          TYPE        CLUSTER-IP      EXTERNAL-IP   PORT(S)    AGE
service/keycloak-discovery    ClusterIP   None            <none>        7800/TCP   5m28s
service/keycloak-service      ClusterIP   172.30.35.121   <none>        8443/TCP   5m28s

NAME                             READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/rhbk-operator    1/1     1            1           13m

NAME                                        DESIRED   CURRENT   READY   AGE
replicaset.apps/rhbk-operator-565b768dcd    1         1         1       13m

NAME                          READY   AGE
statefulset.apps/keycloak     1/1     5m28s

NAME
route.route.openshift.io/keycloak-ingress-vvpv8   ...
```

```
$ oc get secret keycloak-initial-admin -o jsonpath='{.data.password}' | base64 -d
xxxxxxxxx
```

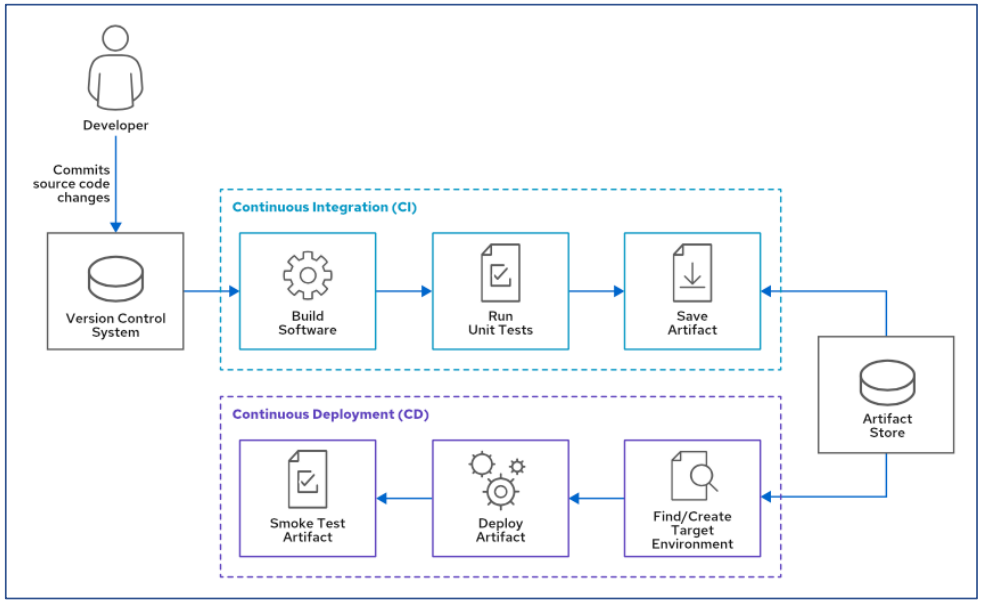


Continuos Integration
Continuos Delivery

→ Developer
→ running application

Jenkins, CruiseControl, TeamCity, GitLab ...
Kubernetes native (Tekton, Argo CD, ...)

GitOps Workflow

→ Administrators
→ live System

Ansible, Puppet, Terraform ...
ArgoCD, FluxCD, JenkinsX

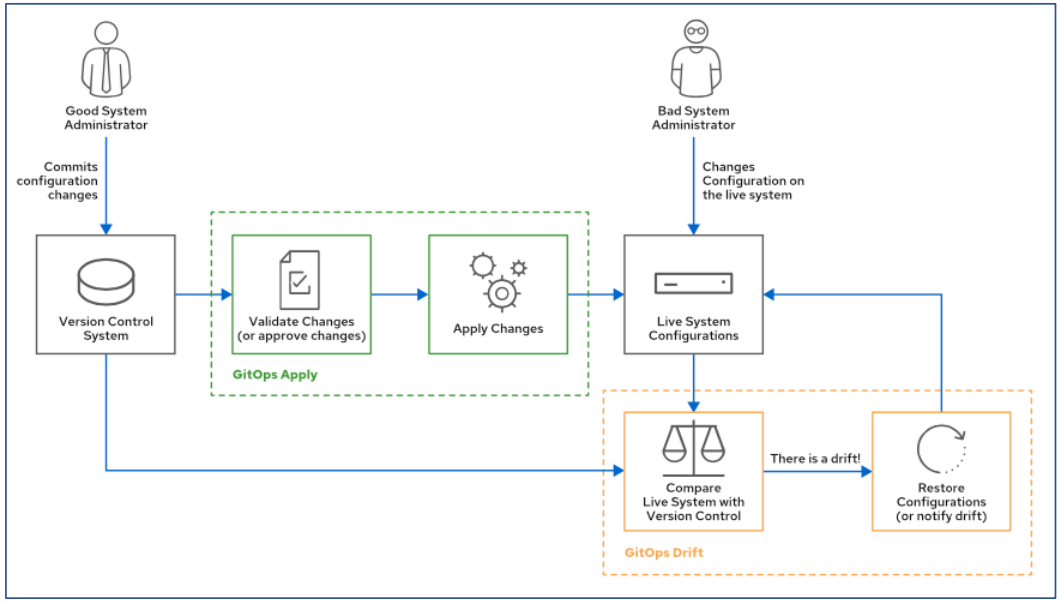

# Tekton - Komponenten

Konfiguration:

- Step
  Script/Programm welches in einem Container ausgeführt werden
  wird innerhalb eines Tasks definiert
- Task
  definieren Ein- und Ausgabeparameter, Umgebung für Steps
  enthalten 1..* Steps
- Pipeline
  definieren Ein- und Ausgabeparameter, Umgebung für Taks
  enthalten 1..* Tasks
- Eventlistener
  reagieren auf HTTP-Events z.B. von VCS

Laufzeit:

- TaskRun
  Ausführung eines Task mit konkreten Parameter, Umgebung (z.B. Workspace)
- PipelineRun
  Ausführung einer Task mit konkreten Parameter, Umgebung (z.B. Workspace)

ClusterTasks : global, vom Operator bereitgestellt

```
$ oc get clustertasks
NAME                        AGE
argocd-task-sync-and-wait   175d
buildah                     175d
git-cli                     175d
git-clone                   175d
...
```

- ➢ kein zentraler Server wie z.B. Jenkins
- ➢ Pipeline, Task / PipelineRun, TaskRun sind *namespaced* Resourcen

# Tekton PipelineRun

- Tasks: können über Ein- und Ausgabeparameter/Bedingungen mit einander verknüpft werden
  die Tasks werden dann ausführt wenn Parameter/Bedingungen erfüllt sind
- Workspace: für alle Tasks sollte ein gemeinsamer Workspace (Persisten-Volume) verwendet werden auf dem gearbeitet wird
  (z.B. ausgechecktes Git-Repo)

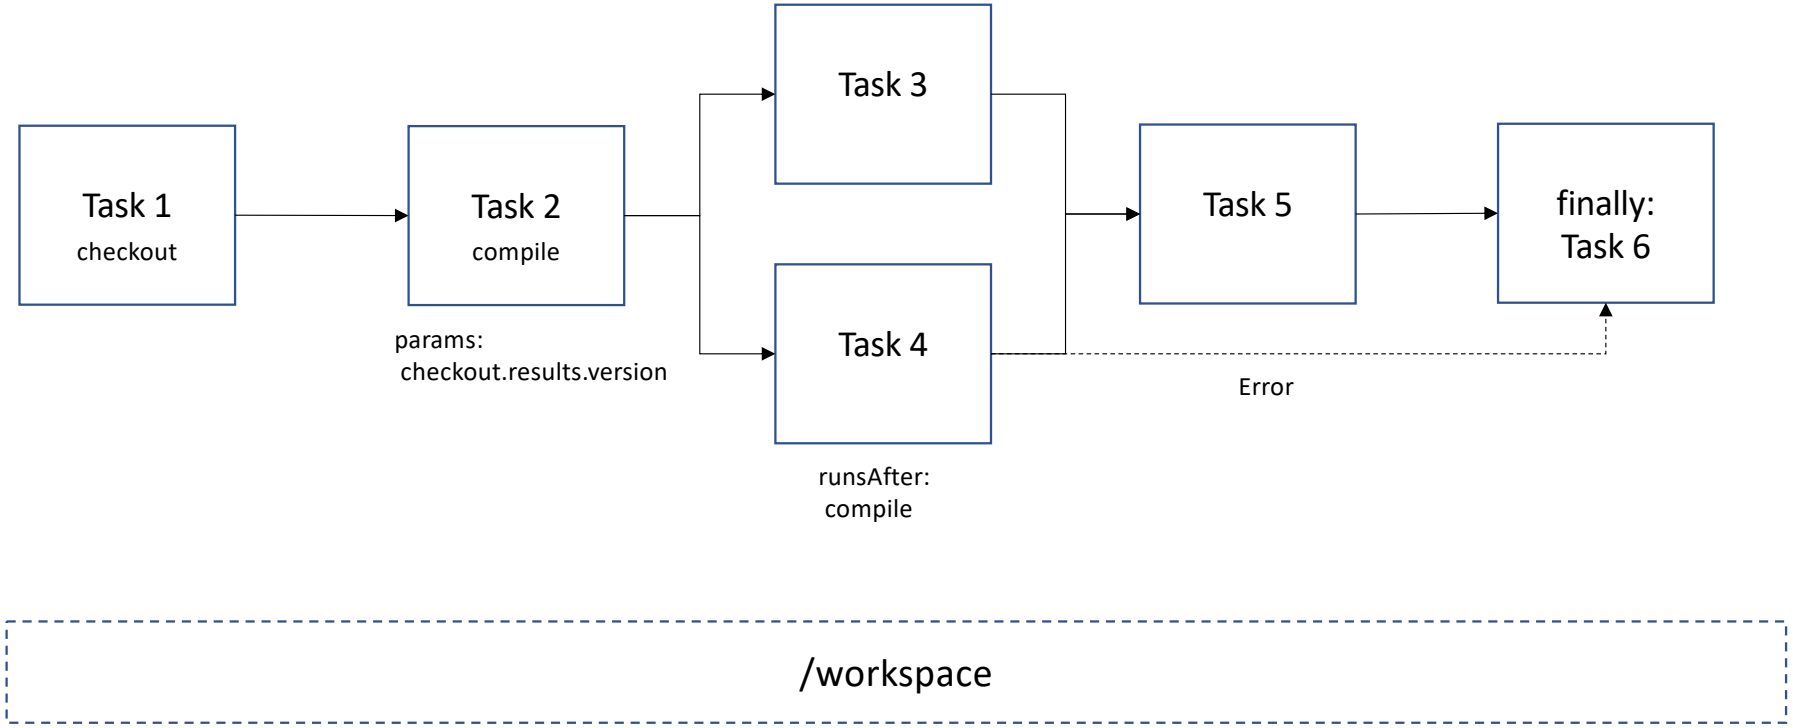

# Tekton Eventlistener

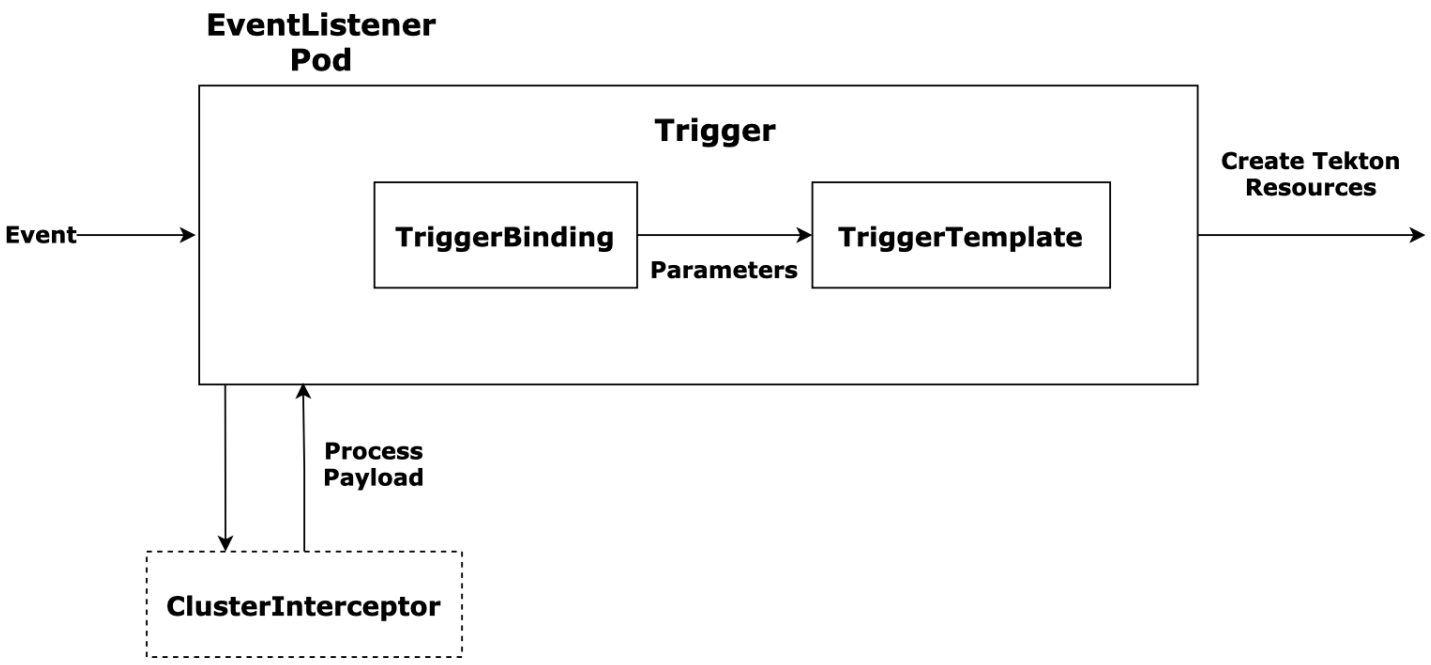

# Tekton Eventlistener

```yaml
apiVersion: triggers.tekton.dev/v1beta1
kind: EventListener
metadata:
  name: webserver-build
spec:
  serviceAccountName: pipeline
  triggers:
  - name: github
    interceptors:
    - ref:
        name: cel
      params:
      - name: filter
        value: body.commits[0].author.name != 'pipeline'
      - name: overlays
        value:
        - key: git_short_rev
          expression: body.head_commit.id.truncate(7)
        - key: git_rev
          expression: body.head_commit.id
    bindings:
    - name: git-rev
      value: $(extensions.git_rev)
    - name: git-short-rev
      value: $(extensions.git_short_rev)
    template:
      ref: webserver-build
```

Deployment
el-webserver-build

Service
el-webserver-build

Route
...
(muss selber erstellt werden)

Tekton Resourcen können normal mit den oc-Client verwendet werden

```
$ oc create -f pr.yml && oc get tr -w

webserver-build-run-qdvnd-fetch-repo       Unknown     Pending
webserver-build-run-qdvnd-build-image      True        Succeeded
webserver-build-run-qdvnd-commit-changes   True        Succeeded
webserver-build-run-qdvnd-fetch-repo       True        Succeeded
```

```
$ oc logs webserver-build-run-qdvnd-fetch-repo-pod
...
{"level":"info","ts":1719519668.1135974,"caller":"git/git.go:1
76","msg":"Successfully cloned git@github.com:dstraub/lab-
sample.git @ ad23aa3566111e6f588f32256bf5e4fdd77b5199
{"level":"info","ts":1719519668.1679106,"caller":"git/git.go:2
15","msg":"Successfully initialized and updated submodules in
path /workspace/output/"}
```

```
$ oc logs el-webserver-build-6d84df79bb-mztzj | tail -n 2 | jq
{
  "severity": "info",
  "timestamp": "2024-06-27T20:29:08.529Z",
  "logger": "eventlistener.event-broadcaster",
  "caller": "record/event.go:298",
  "message": "Event(v1.ObjectReference{Kind:\"EventListener\", Namespace:\"webserver-build\", Name:\"webserver-build\",
UID:\"65993abf-ccf6-4aa4-a9ca-957bbda2e2a7\", APIVersion:\"triggers.tekton.dev/v1beta1\", ResourceVersion:\"22473585\",
FieldPath:\"\"}): type: 'Normal' reason: 'dev.tekton.event.triggers.done.v1' ",
  "commit": "f76be74"
}
{
  "severity": "info",
  "timestamp": "2024-06-27T20:29:08.540Z",
  "logger": "eventlistener",
  "caller": "sink/sink.go:420",
  "message": "interceptor stopped trigger processing: rpc error: code = FailedPrecondition desc = expression
body.commits[0].author.name != 'pipeline' did not return true",
  "commit": "f76be74",
  "eventlistener": "webserver-build",
"/triggers-eventid": "2fcf657c-f678-4189-a218-06d45ddfba12",
  "eventlistenerUID": "65993abf-ccf6-4aa4-a9ca-957bbda2e2a7",
  "/trigger": "github"
}
```

GitOps – Workflow mit Pipelines:

- Apply Pipeline:
  - `validate : oc apply --validate --dry-run [ folder/files from Git ]`
  - `apply    : oc apply`

- Drift Pipeline:
  - `diff     : oc diff [ folder/files from Git ]`
  - `optional/restore: oc apply`

GitOps – Workflow mit ArgoCD / FluxCD:
Abgleich Ist-Zustand (Cluster) mit Kustomize/Helm-Definitionen im Git
Benachrichtigungen, manueller/automatische Synchronisation bei Abweichungen

| | | | | |
|---|---|---|---|---|
| apps<br>calibre | ssh://git@gitea.apps:10022/ds/calibre.git/overlays/production<br>in-cluster/apps | HEAD | Healthy<br>Synced | ⋮ |
| apps<br>pgadmin | ssh://git@gitea.apps:10022/ds/pgadmin.git/overlays/production<br>in-cluster/apps | HEAD | Healthy<br>Synced | ⋮ |
| apps<br>postgres | ssh://git@gitea.apps:10022/ds/postgres.git/overlays/production<br>in-cluster/database | HEAD | Healthy<br>Synced | ⋮ |
| apps<br>rest-sample | ssh://git@gitea.apps:10022/ds/rest-sample.git/overlays/production<br>in-cluster/sample | HEAD | Healthy<br>OutOfSync | ⋮ |

Red Hat Openshift GitOps - Operator

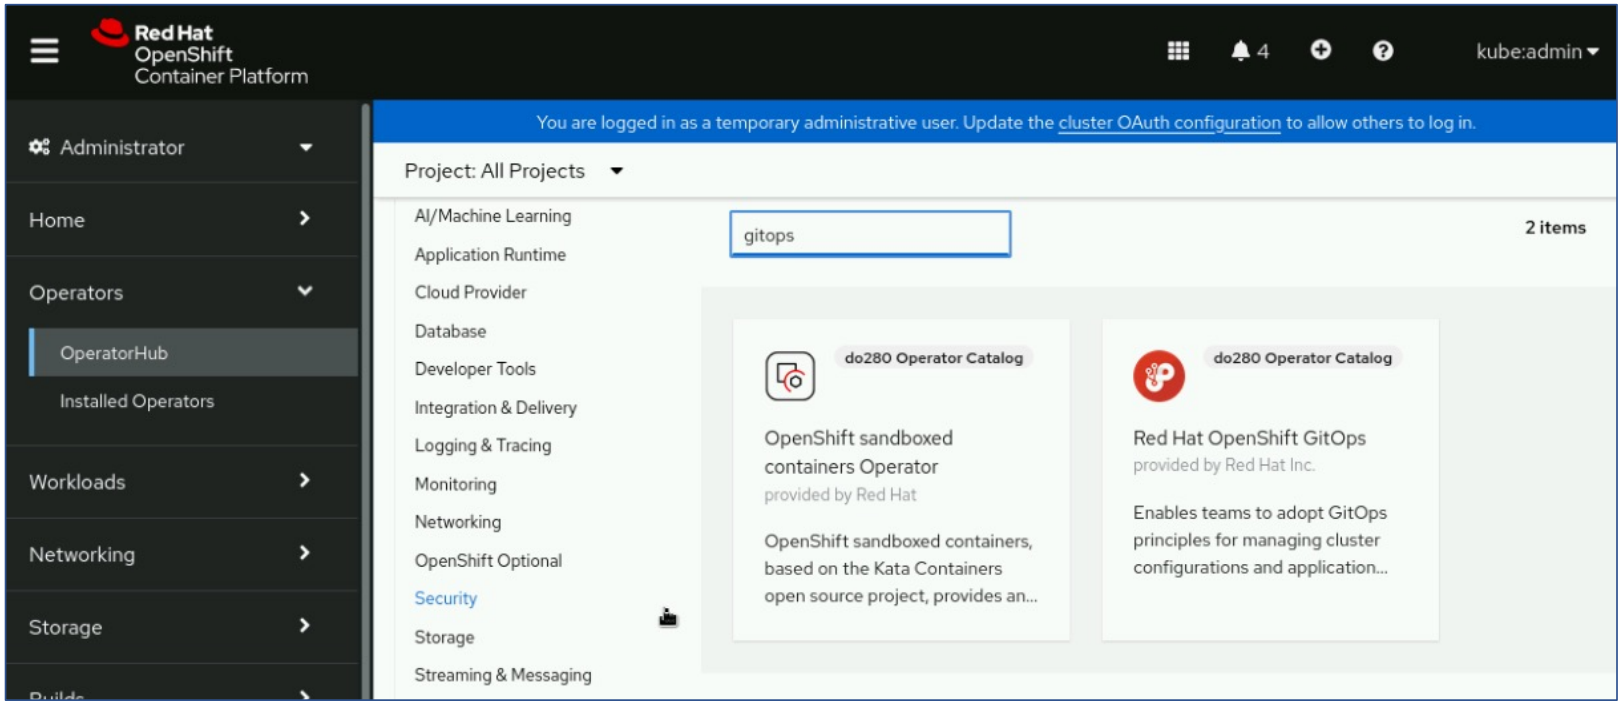


ArgoCD = Openshift GitOps

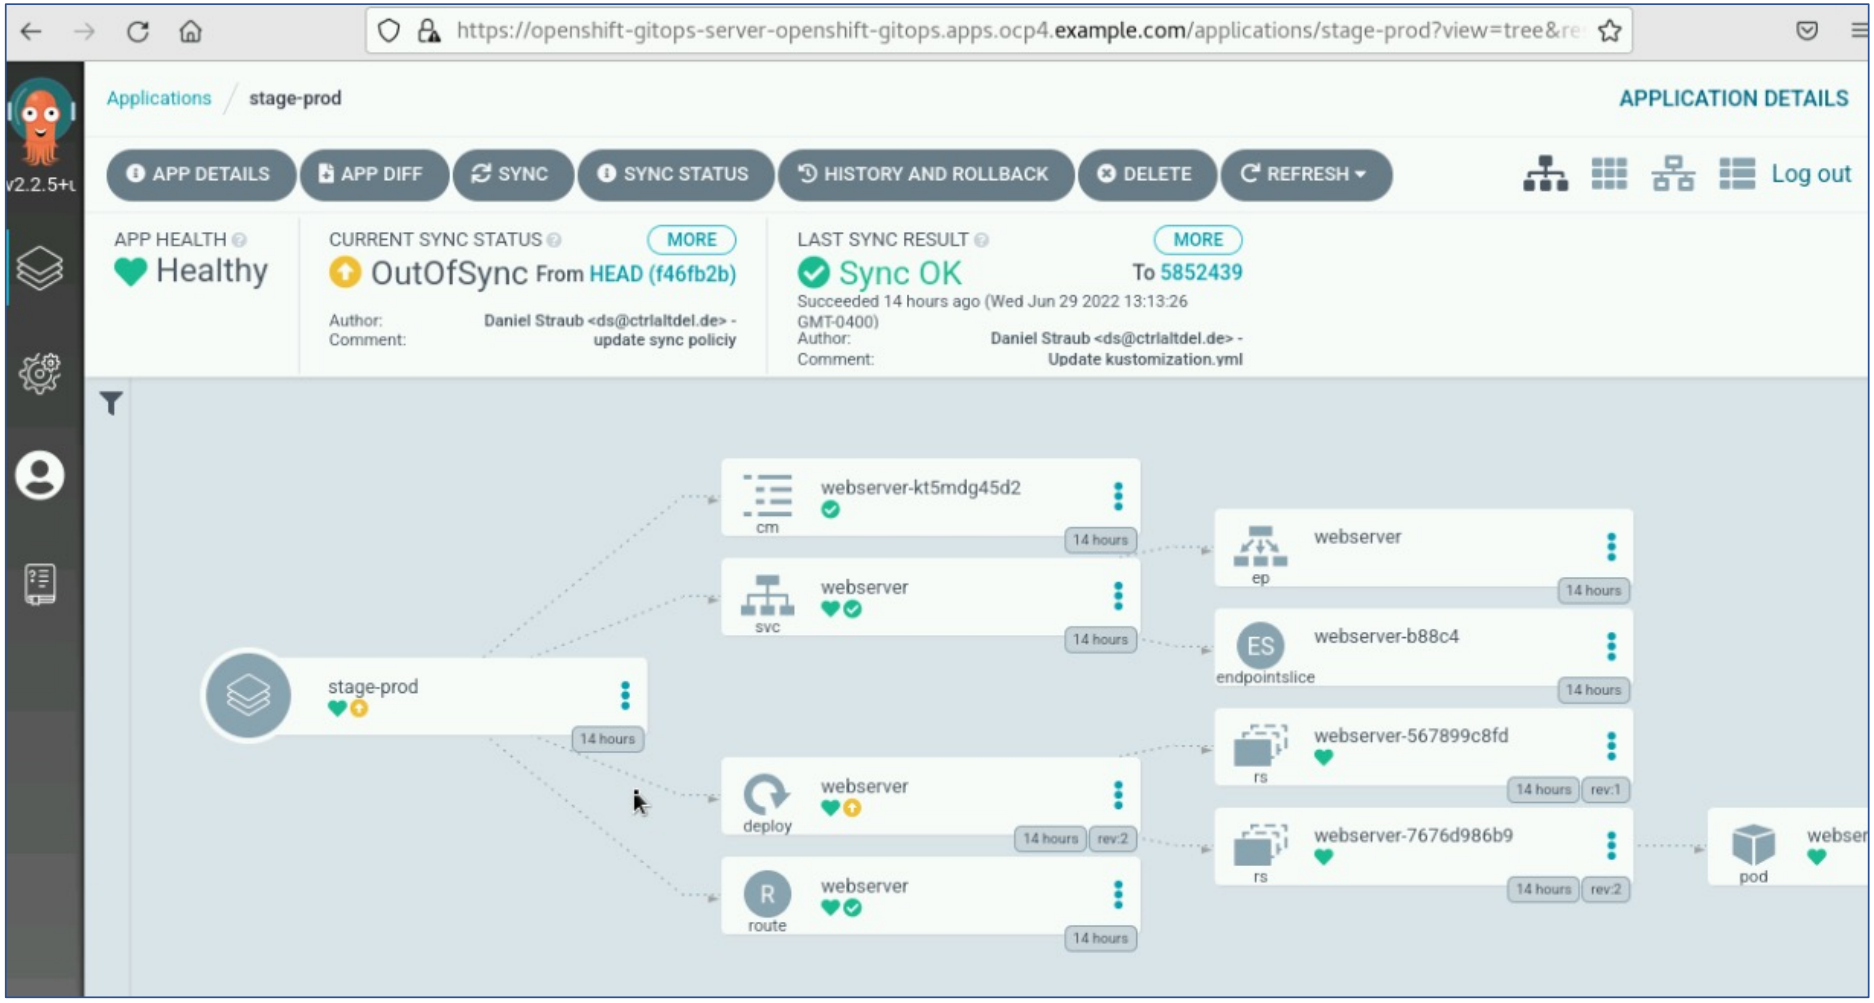


Openshift GitOps